

# Introduction to High Performance Computing (HPC)

**sciCORE**  
Center for scientific computing

**Swiss TPH**   
Swiss Tropical and Public Health Institute  
Schweizerisches Tropen- und Public Health-Institut  
Institut Tropical et de Santé Publique Suisse

  
**Swiss Institute of  
Bioinformatics**

  
**Universität  
Basel**

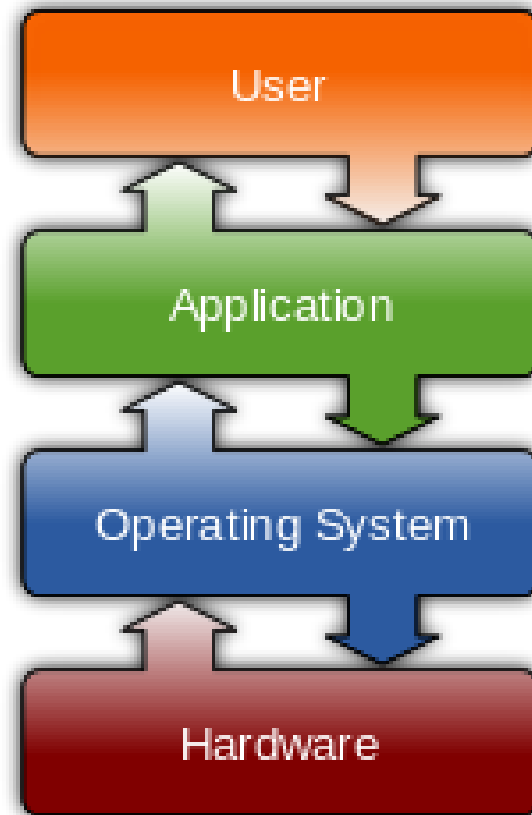
# Outline



- 09:30 – Course starts
- 09:30 – 11:00: Introduction to Linux. Part 1
- 11:00 – 11:15: Coffee break
- 11:15 – 12:30: Introduction to Linux. Part 2
- 12:30 – 14:00: Lunch
- 14:30 – 16:00: Cluster usage. Part 1
- 16:00 – 16:15: Coffee break
- 16:15 – 17:30: Cluster usage. Part 2
- 17:30 – 18:00: Feedback + Q&A

I will assume no previous knowledge on Linux or on clusters usage.

# What's an OS?



# What's an OS?



Sun



HP



IBM



VMWare



Apple

ORACLE

Oracle



OS X



Windows



Linux



Xen



Red Hat



Fedora



CentOS



Debian



Ubuntu



Mint



SUSE



Mageia



Arch Linux



Slackware



Mandriva



Gentoo



FreeBSD



OpenBSD



NetBSD

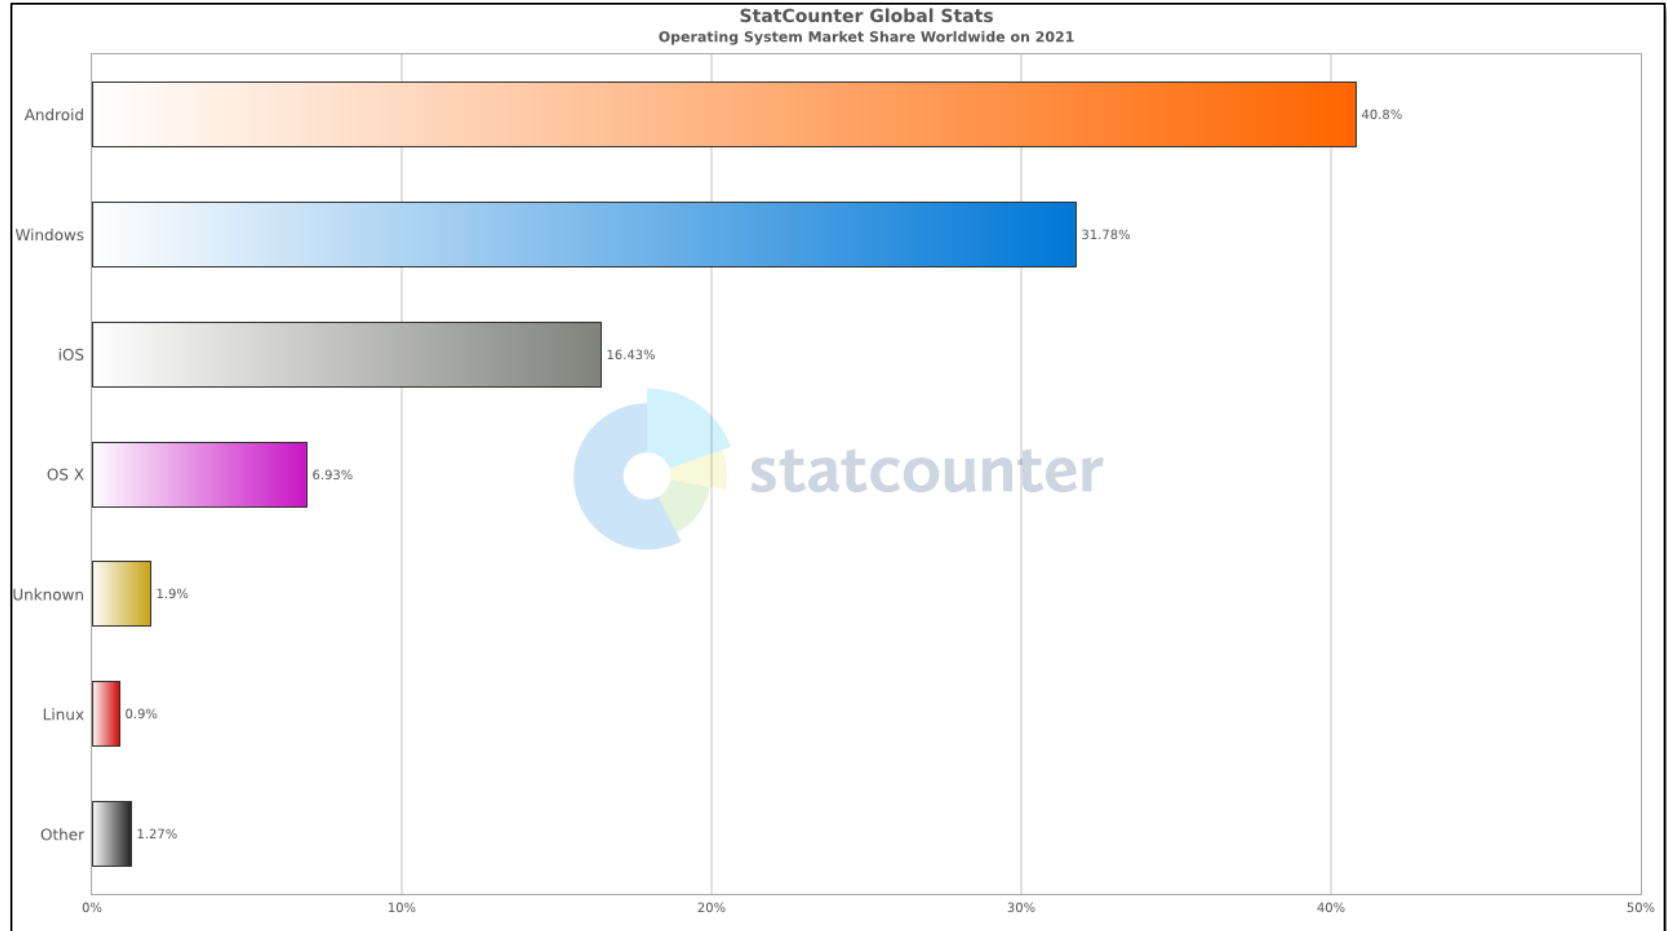


DragonFly BSD

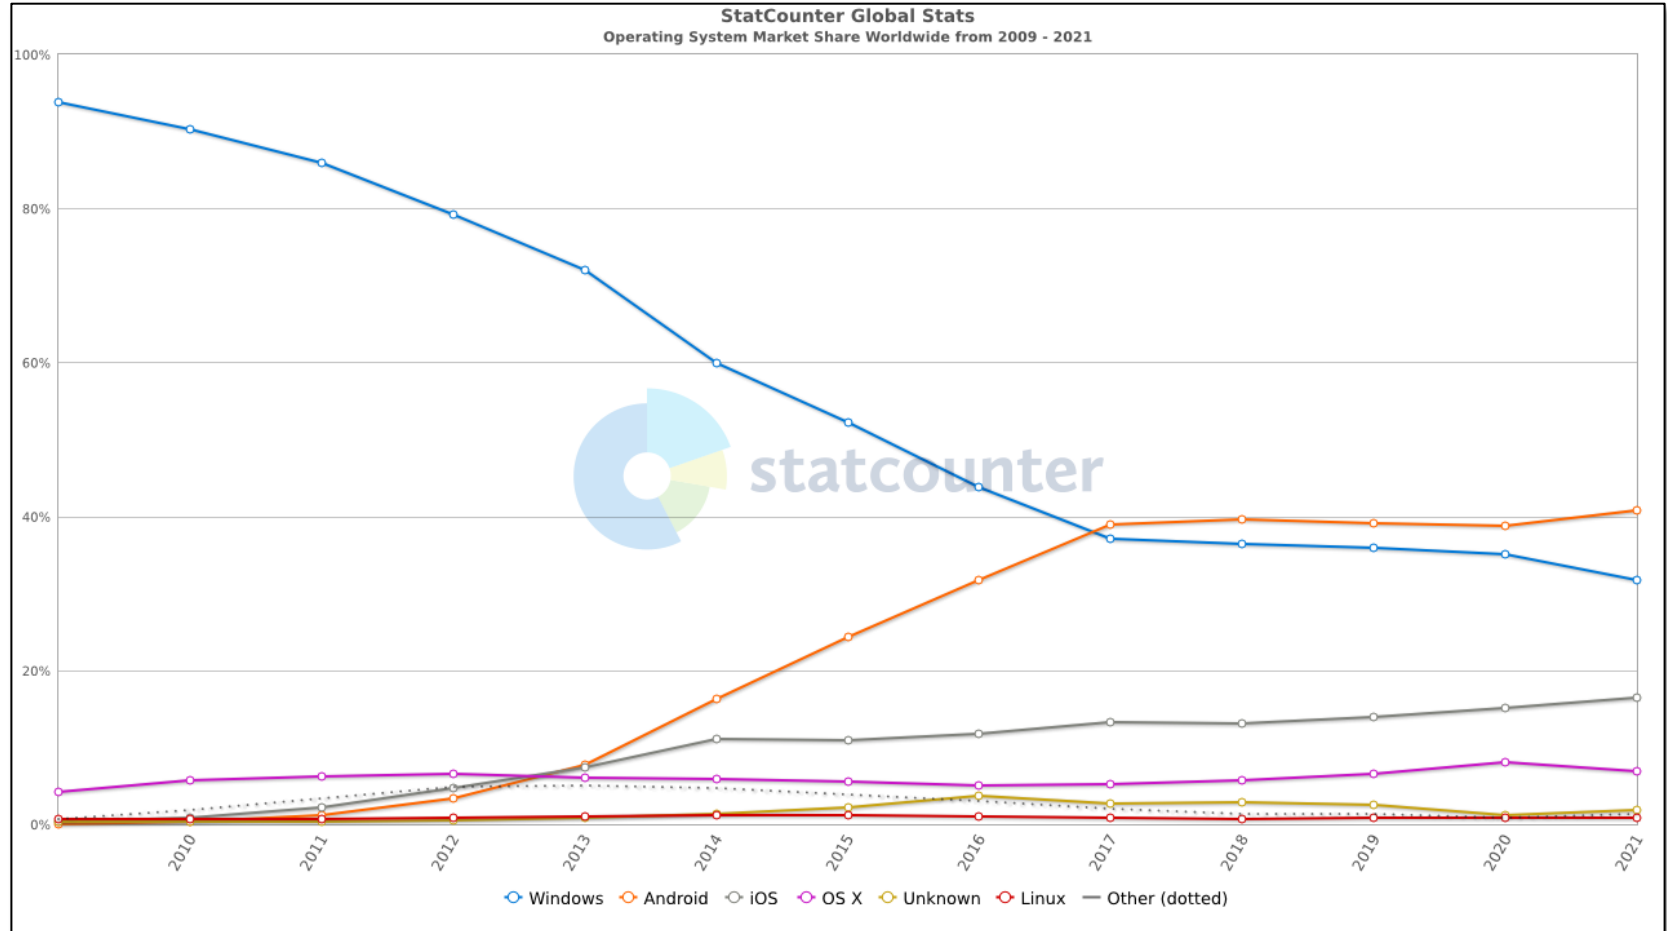


Darwin

# What's an OS?

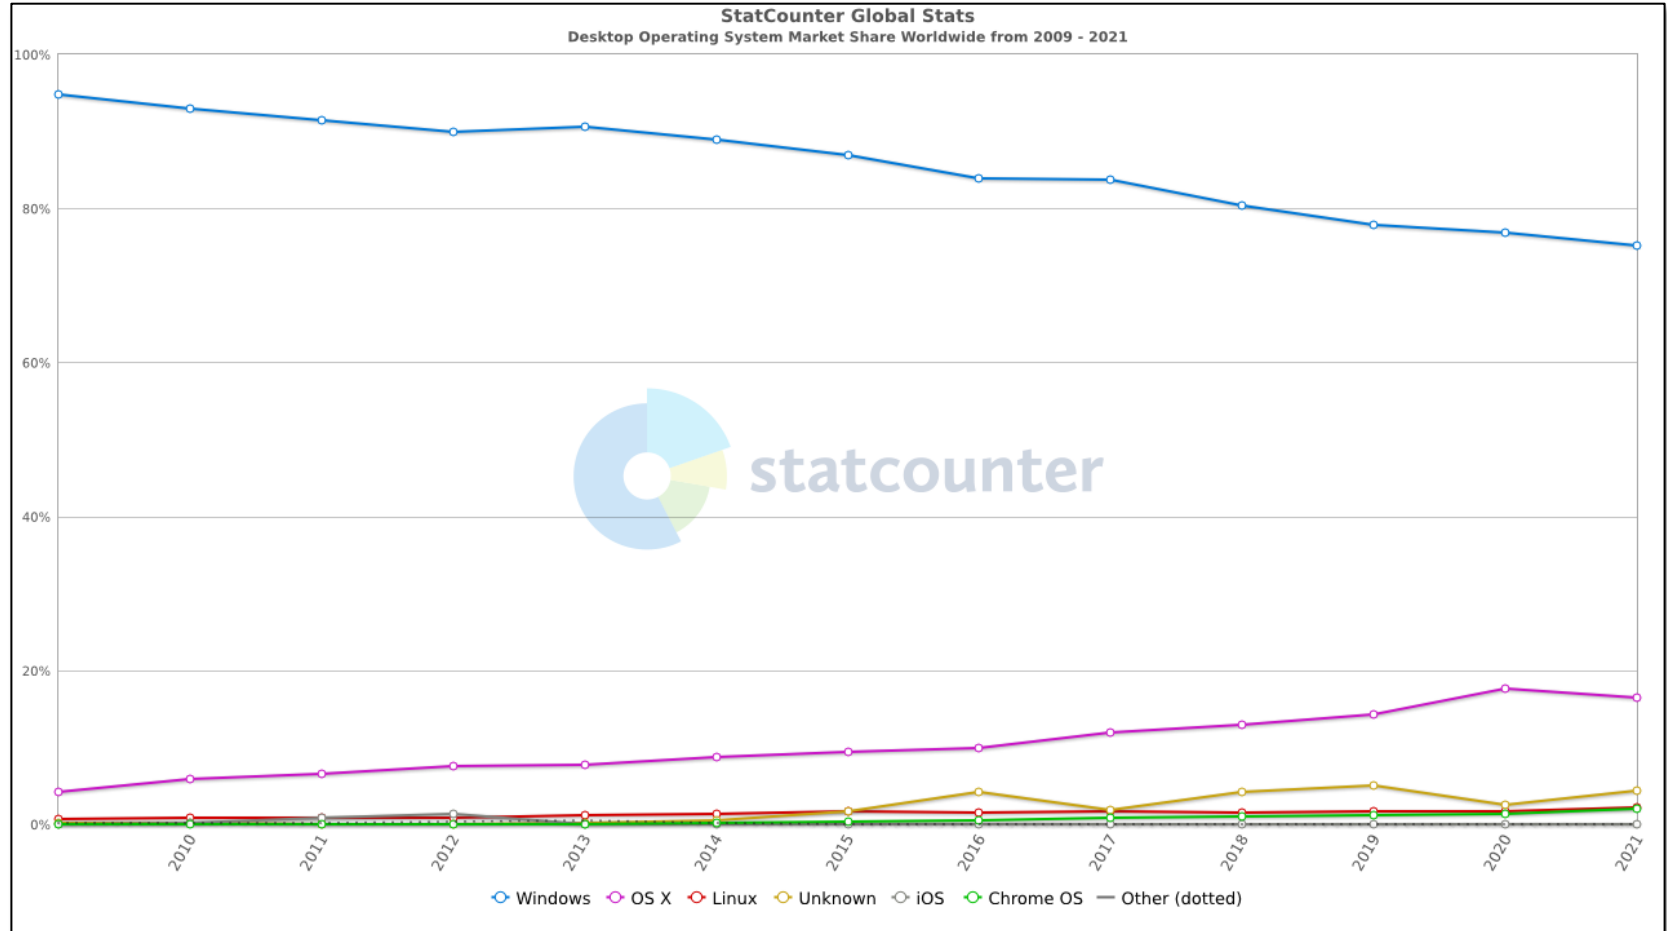


# What's an OS?



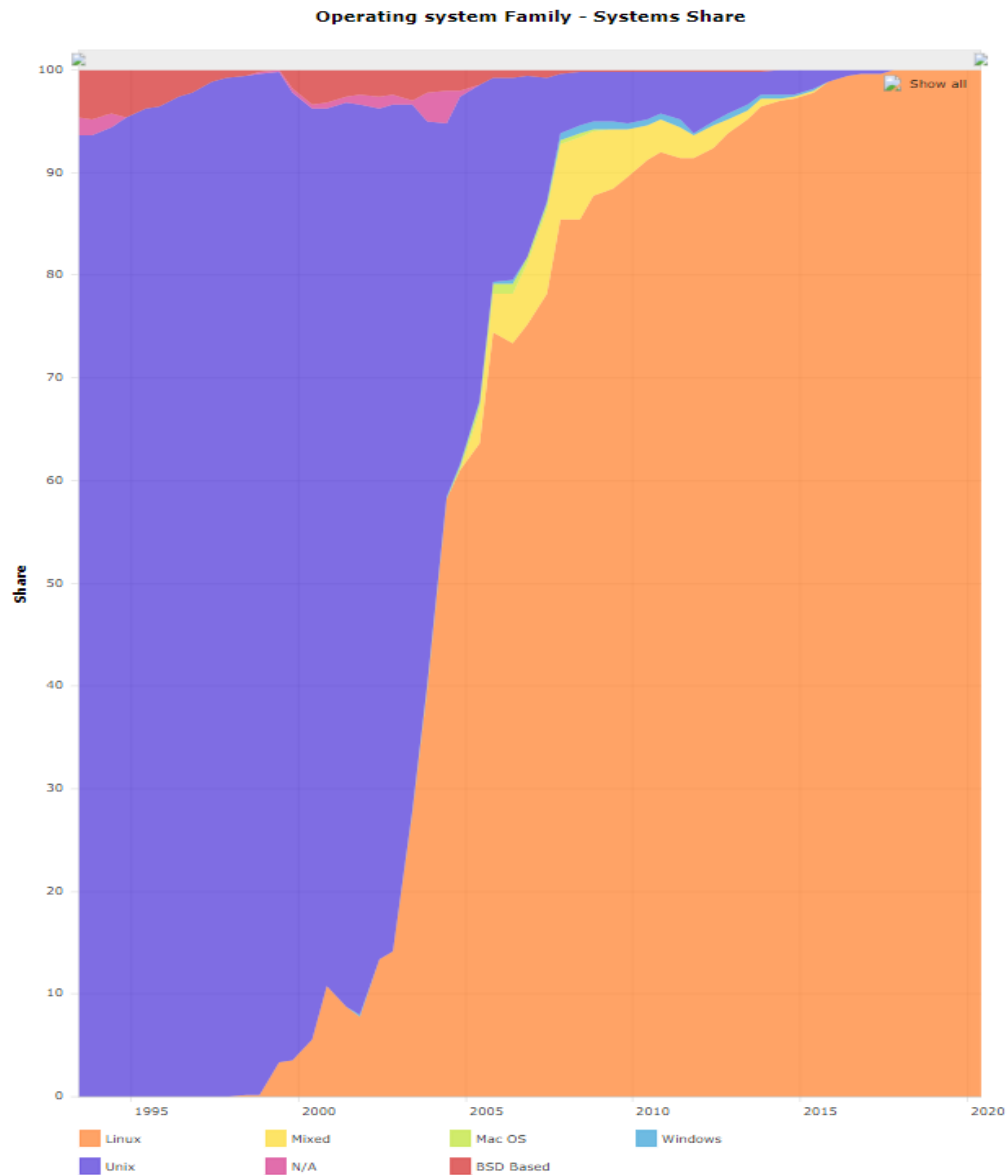
Source: statcounter.com

# What's an OS?



Source: statcounter.com

# What's an OS?

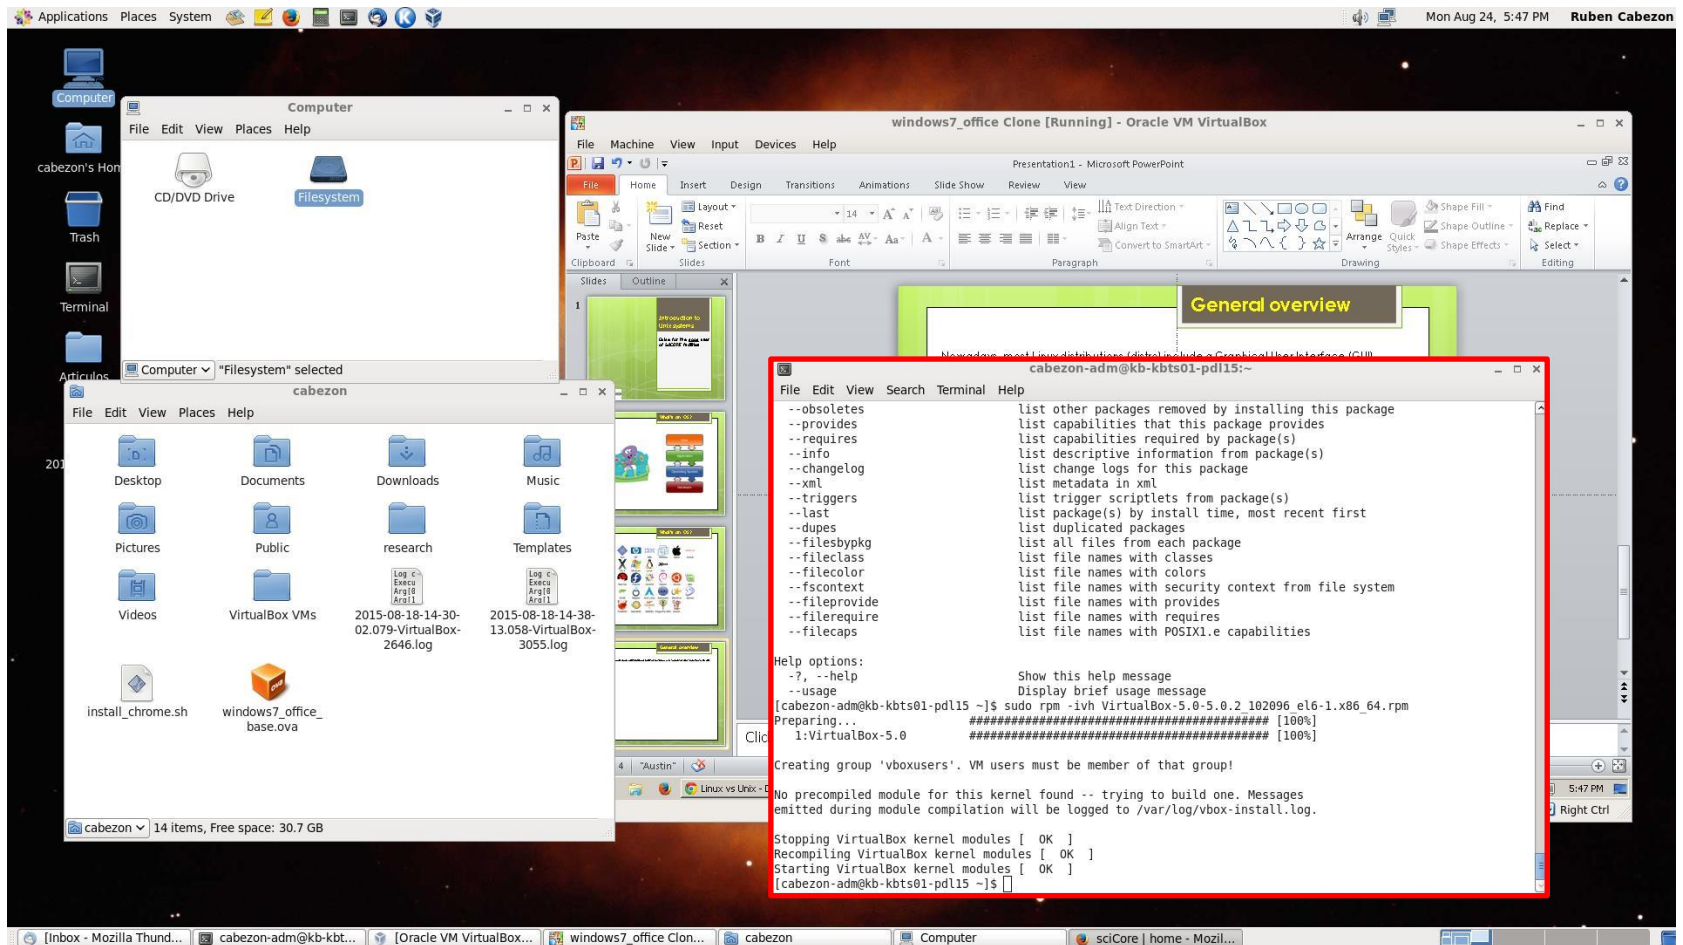


Source: top500.org



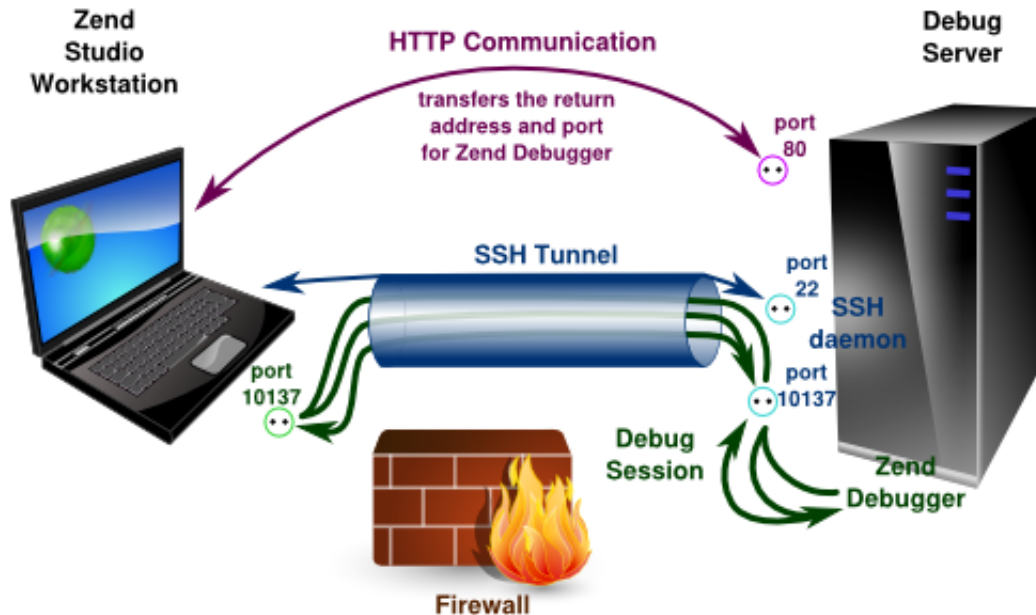
# General overview

Nowadays, most Linux distributions (distro) include a Graphical User Interface (GUI).  
Usual ones: KDE, GNOME



# General overview

When using clusters or supercomputers we log in remotely, usually via a secure connection: **ssh**



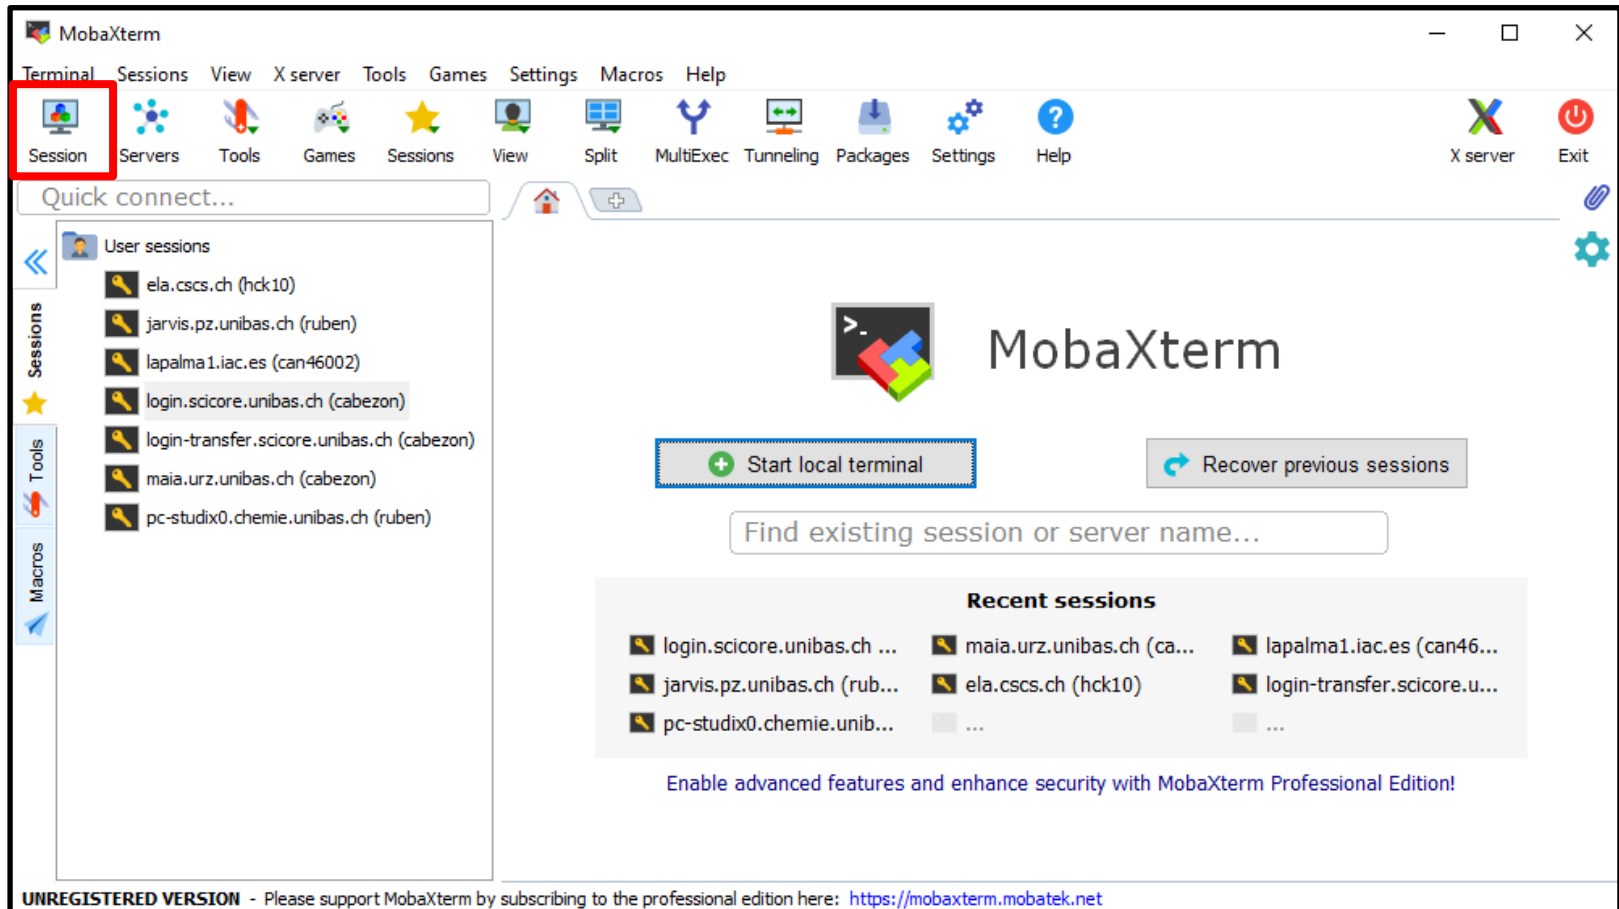
This will open a terminal, in which we will log in and work.

**NOTE:** Windows users without ssh, download MobaXterm  
<http://mobaxterm.mobatek.net/download-home-edition.html>

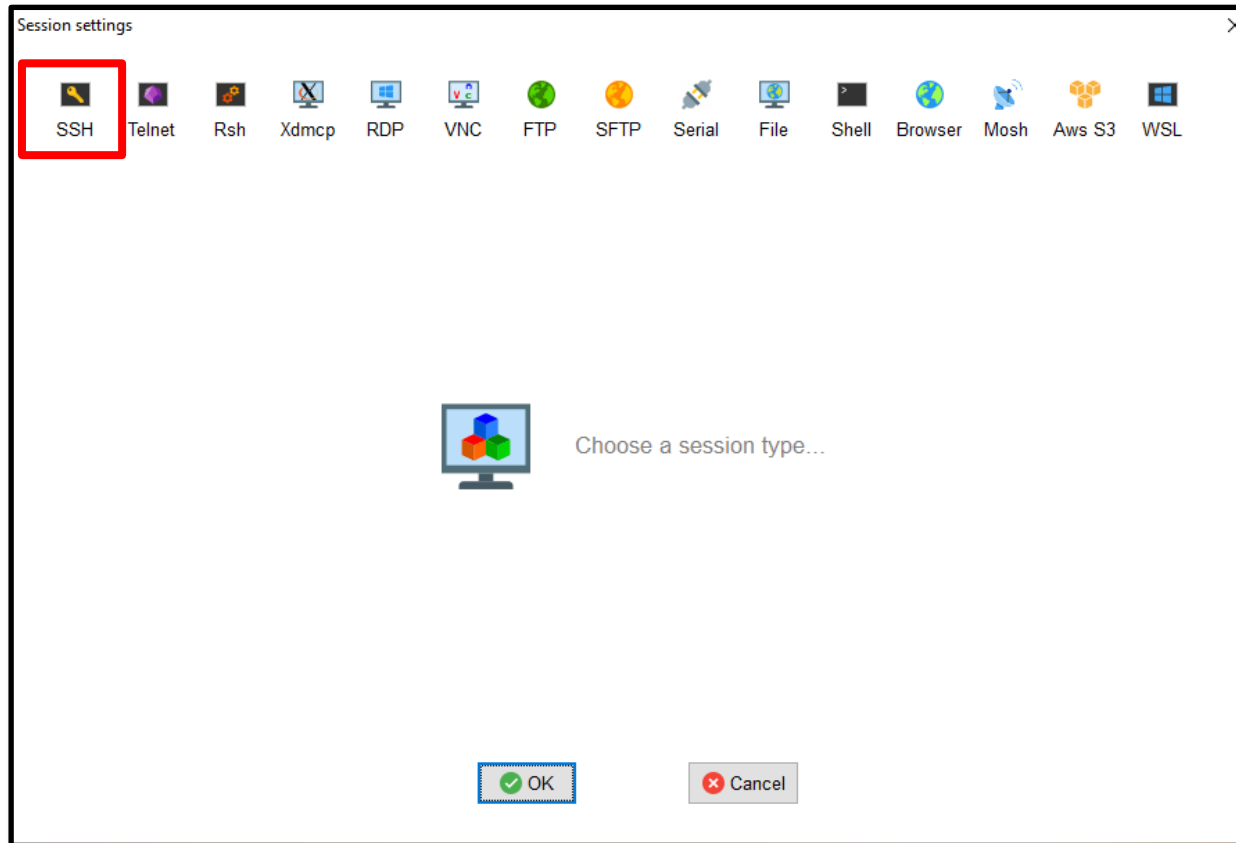
**NOTE:** Mac users without X11, download XQuartz  
<http://xquartz.macosforge.org>

Another option is PuTTY client + Xming  
<http://the.earth.li/~sgtatham/putty/latest/x86/putty.exe>  
<http://sourceforge.net/projects/xming/>

# Using MobaXterm



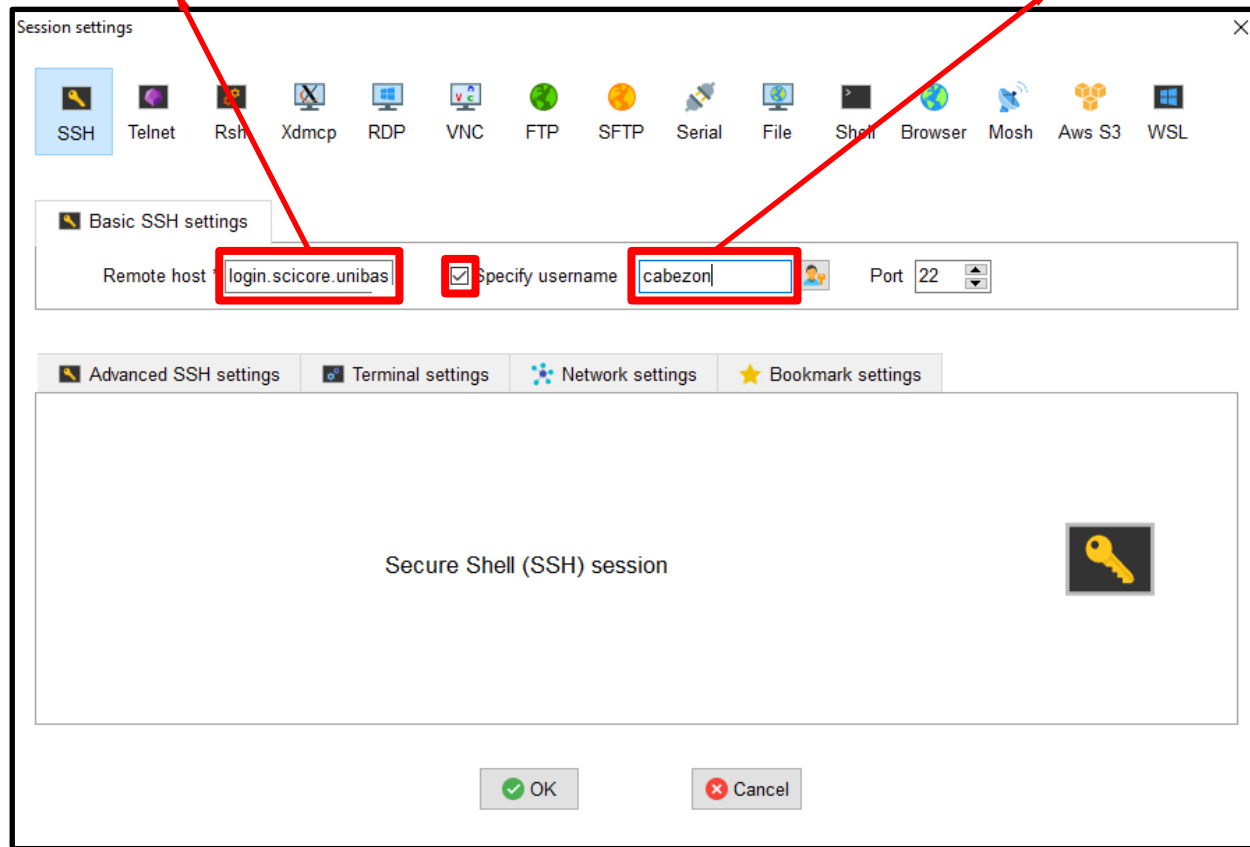
# Using MobaXterm



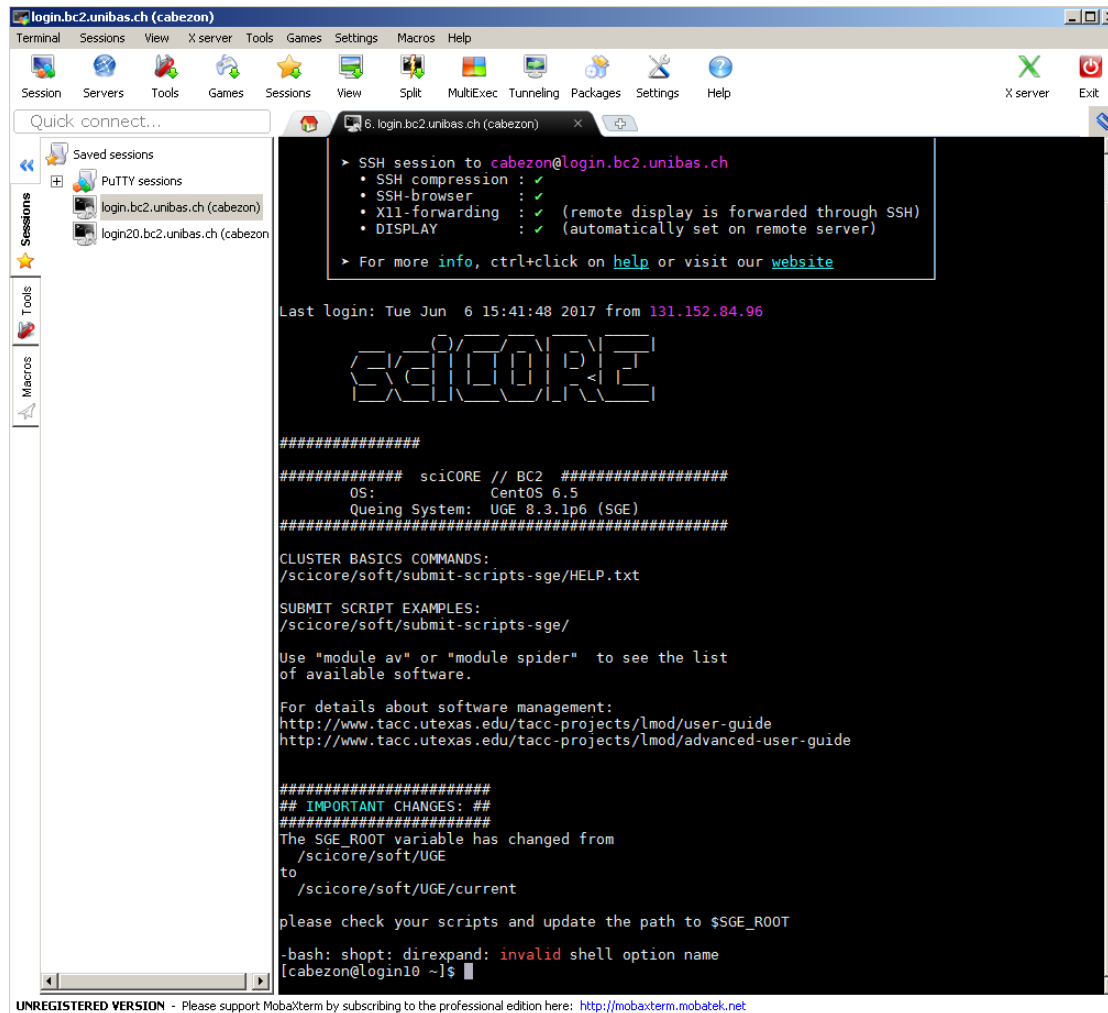
# Using MobaXterm

86.119.35.191

<your username>



# Using MobaXterm



The screenshot shows the MobaXterm application window with the title bar "login.bc2.unibas.ch (cabezon)". The menu bar includes Terminal, Sessions, View, X server, Tools, Games, Settings, Macros, and Help. The toolbar contains icons for Session, Servers, Tools, Games, Sessions, View, Split, MultiExec, Tunneling, Packages, Settings, Help, X server, and Exit. The left sidebar shows "Quick connect...", "Saved sessions", "PUTTY sessions", "Sessions" (with a list of login sessions), "Tools", and "Macros". The main terminal area displays the following text:

```
> SSH session to cabezon@login.bc2.unibas.ch
• SSH compression : ✓
• SSH-browser      : ✓
• X11-forwarding   : ✓ (remote display is forwarded through SSH)
• DISPLAY          : ✓ (automatically set on remote server)

> For more info, ctrl+click on help or visit our website

Last login: Tue Jun  6 15:41:48 2017 from 131.152.84.96

  sciCORE

#####
##### sciCORE // BC2 #####
OS:          CentOS 6.5
Queuing System: UGE 8.3.lp6 (SGE)
#####

CLUSTER BASICS COMMANDS:
/scicore/soft/submit-scripts-sge/HELP.txt

SUBMIT SCRIPT EXAMPLES:
/scicore/soft/submit-scripts-sge/

Use "module av" or "module spider" to see the list
of available software.

For details about software management:
http://www.tacc.utexas.edu/tacc-projects/lmod/user-guide
http://www.tacc.utexas.edu/tacc-projects/lmod/advanced-user-guide

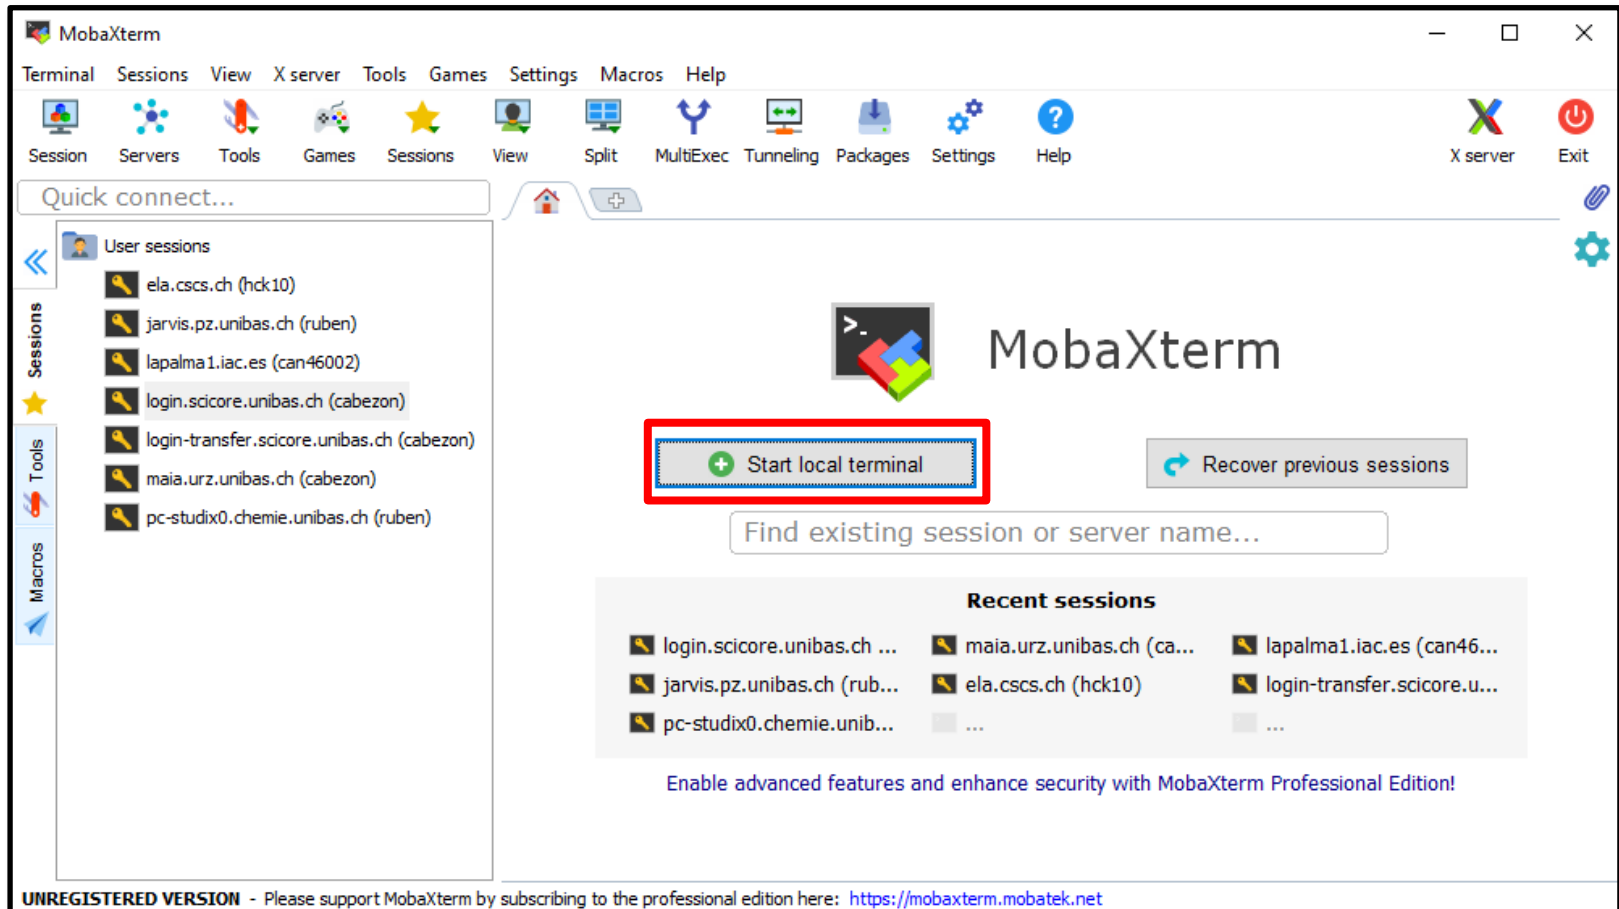
#####
## IMPORTANT CHANGES: ##
#####
The SGE_ROOT variable has changed from
/scicore/soft/UGE
to
/scicore/soft/UGE/current
please check your scripts and update the path to $SGE_ROOT

-bash: shopt: direxpand: invalid shell option name
[cabezon@login10 ~]$
```

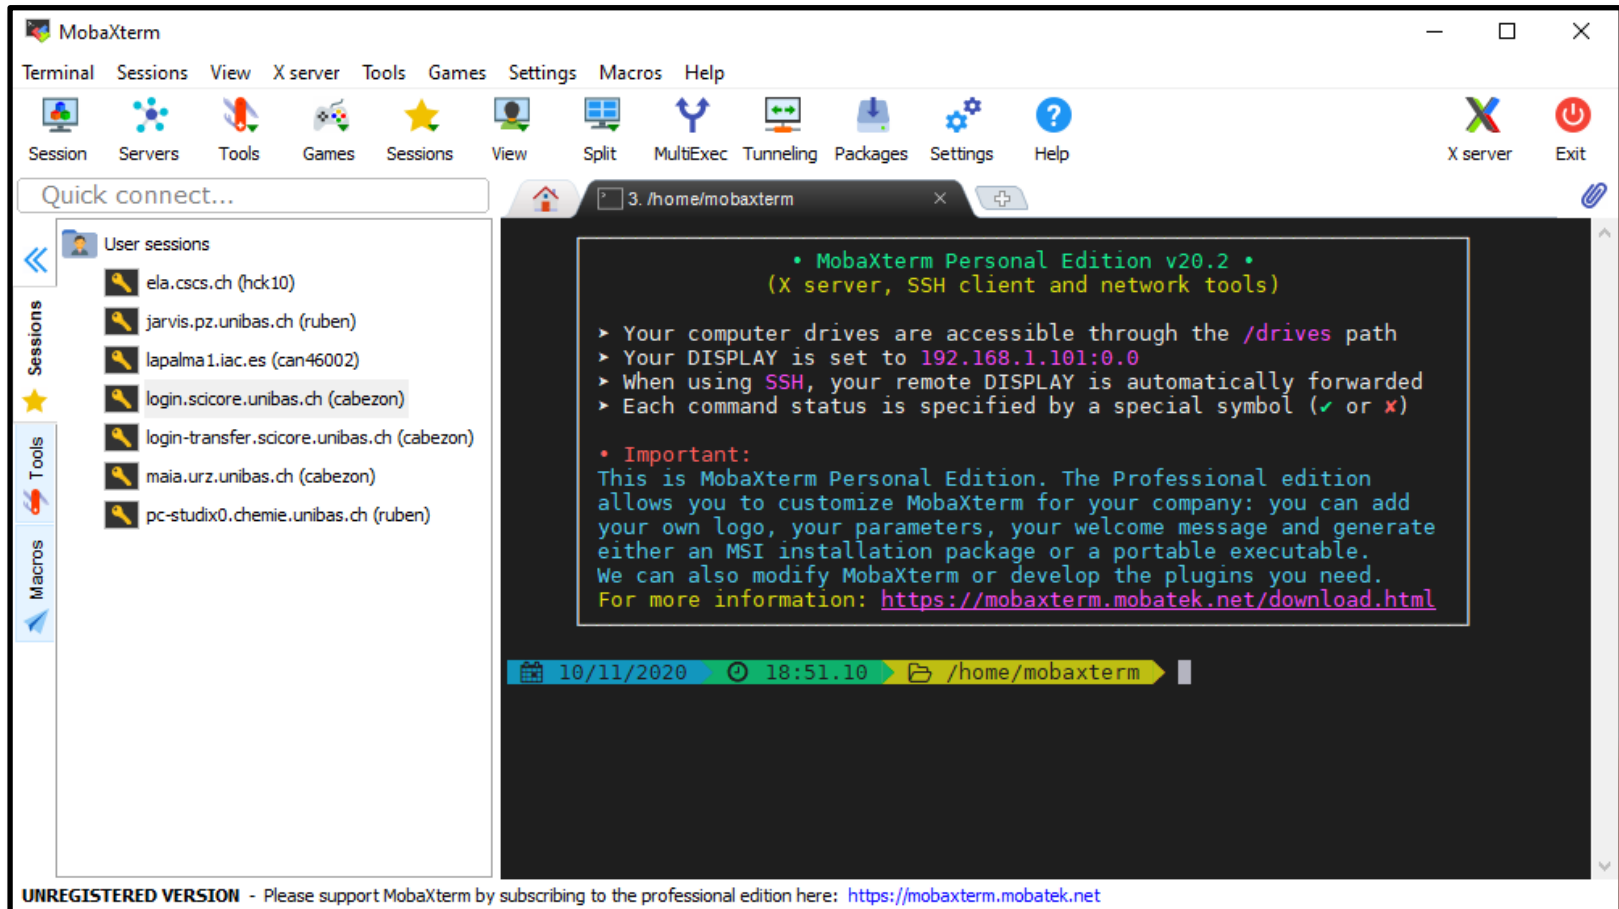
At the bottom of the window, a message reads: "UNREGISTERED VERSION - Please support MobaXterm by subscribing to the professional edition here: <http://mobaxterm.mobatek.net>".

ssh [options] <username>@86.119.35.191

# Using MobaXterm

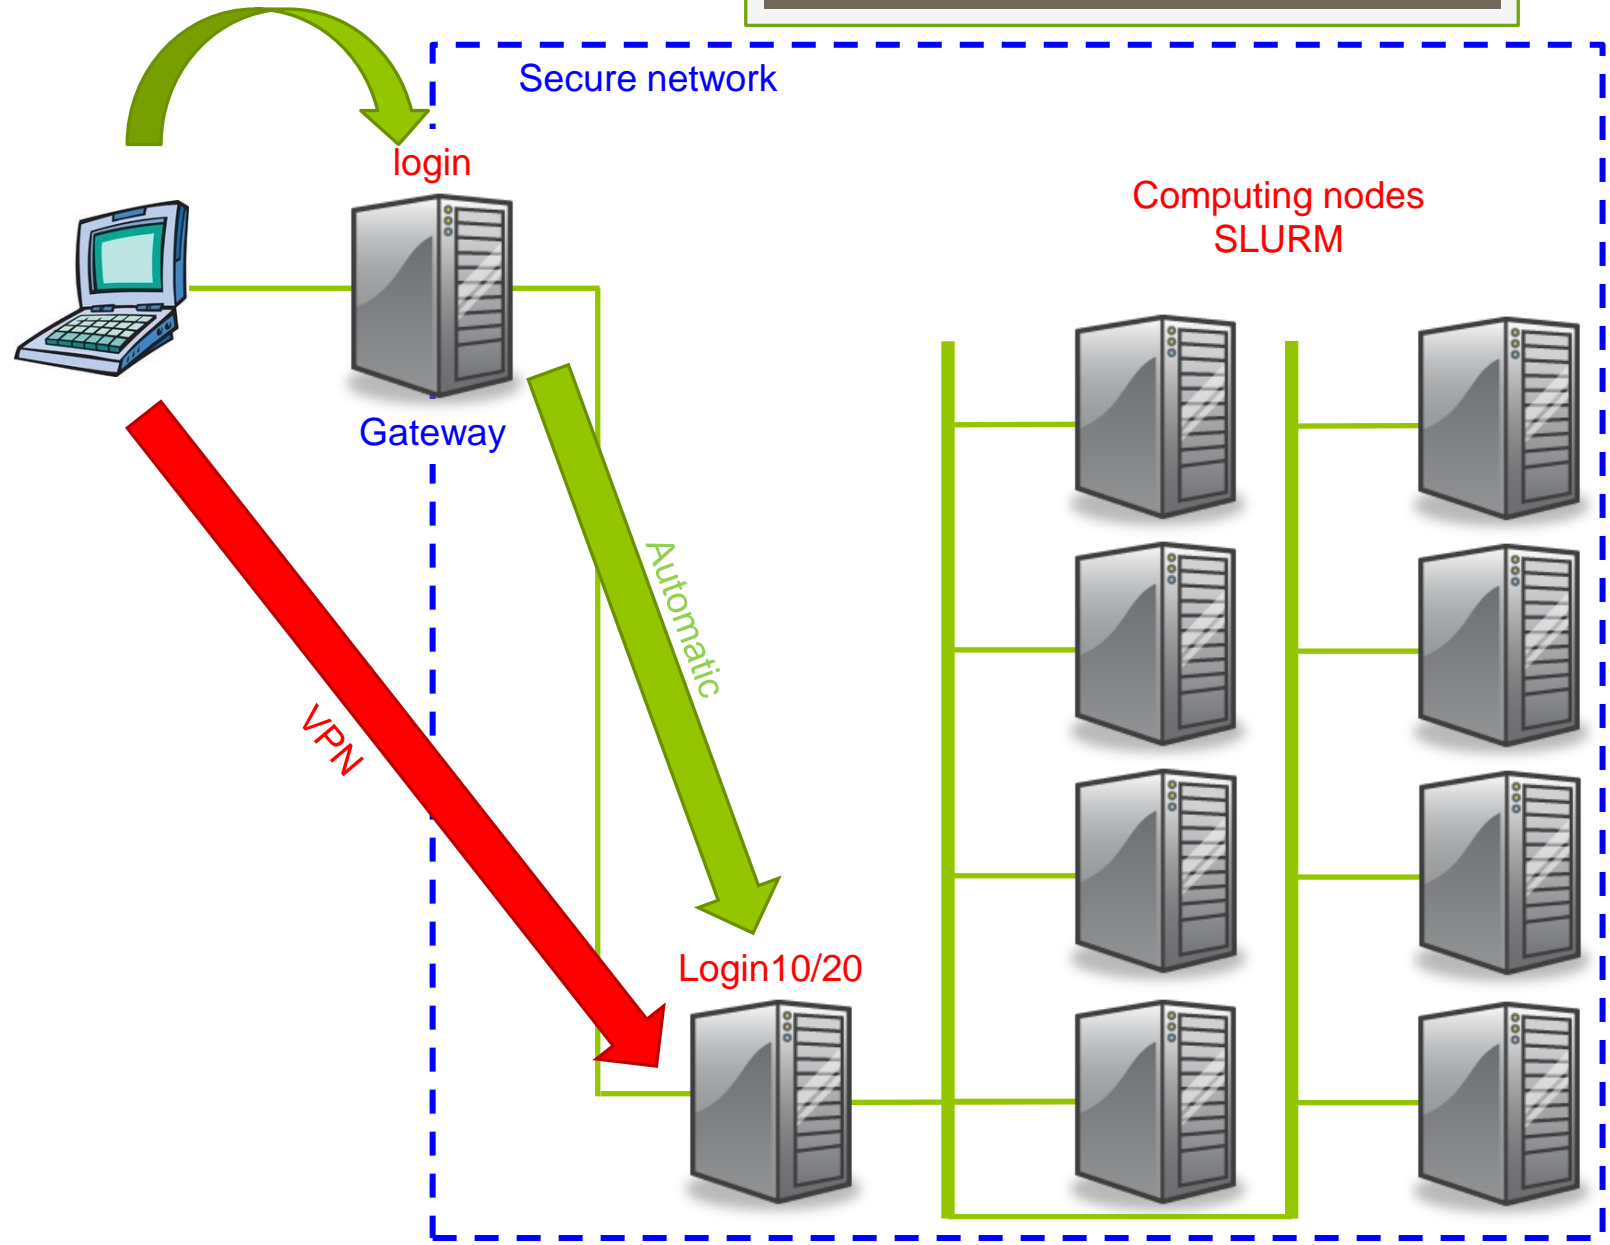


# Using MobaXterm





# Using MobaXterm



# Nomenclature

Something in **red** means that it has to be written literally in the command line.  
Something in **blue** is not a literal command but information for you.

Something between < > has to be completed correspondingly.

Ex: **queue -u <username>**  **queue -u ruben**

Ex: Find information about the command **queue**

**ATTENTION:** Capital letters, blank spaces, dots (.), hyphens (—), and slash(/)

**Linux is case-sensitive!**

## 0. Exercise

1. Log in to the cluster (`ssh -Y <username>@86.119.35.191`)  
(If you have windows, simply use MobaXterm to connect)
2. Check that the computer knows who you are with `who am i`
3. Clear the screen with `clear`
4. Write `yes`. Stop the command with `CTRL+C`
5. Write `sleep 5` and hit Enter. What happens?
6. Write `time sleep 5` and hit Enter. What information does `time` provide?
7. Close the connection with `CTRL+D`.

Do you already know all this?

1. Do the exercises of the handout.
2. Write a script that prints in the screen the total size (in Mb) of the contents of a folder passed as an argument, without using `du`  
**Hint:** You might need to use `ls`, `awk`, `echo` and `bc`. Google a bit to learn about these if you don't know them already.

## 0. Exercise

### Summary

Command	Used to...	Common options
ssh	establish a secure connection <b>ssh [options] &lt;username&gt;@hostname</b>	-X (enables trusted X11 forwarding) -Y (enables untrusted X11 forwarding)
who am I	show user connection info	
clear	clear the screen	
logout	log out of the session	
yes	continuously output 'y'	
sleep	puts the CPU to sleep <b>sleep &lt;number of seconds&gt;</b>	
time	provides the cpu, user and wallclock time of any command	<b>/usr/bin/time -v &lt;command&gt;</b>
CTRL + C	interrupts a running command	
CTRL + D	shortcut to logout	

# 1. Exercise

1. Log in to the cluster
2. List the contents of the temporary directory: `ls /tmp`
3. Learn about `ls` by looking at its manual: `man ls` (press `q` to exit). Try also `info ls` (press `q` to exit).
4. Enter again in the manual of `ls` and search for the option `-1` by writing: `/-1` You can search the next one forwards pressing `n` and backwards with `SHIFT+n`
5. Try more options of `ls` (`-l`, `-lt`, `-lS`, `-m`, `-1`, `-lh`, etc...)
6. Create a directory named `test`: `mkdir test`
7. Check that the directory exists: `ls`
8. Enter in `test/`: `cd test`
9. Exit `test/`: `cd ..` (attention to the blank space after 'cd' !)
10. Create a file containing the list of files of the `tmp` directory: `ls /tmp > data.d`
11. Check that the file `data.d` exists
12. Check the contents of `data.d`: `less data.d` (press `q` to exit)
13. Move `data.d` into `test/`: `mv data.d test` (do you need `/` ?)
14. Enter in `test/` and check that `data.d` is in there.
15. Copy `data.d` into a file named `data2.d`: `cp data.d data2.d`
16. Delete `data.d`: `rm data.d` (careful!!)
17. With the commands that you know figure out a way of renaming `data2.d` to `finaldata.d`
18. Create a directory named `Exercise1` and move `finaldata.d` inside
19. Go up one level and rename `test/` to `linux_course/`

The final outcome should be a directory structure `linux_course/Exercise1` that includes a file named `finaldata.d` with the list of files from `tmp/`

# 1. Exercise

## Summary

Command	Used to...	Common options
ls	list files in a directory	-l (long listing), -h (human readable), -S (sort by size), -t (sort by modification time) -1 (show 1 file / line)
cd	change directory <b>cd &lt;path&gt;</b>	
.	point to the present directory	
..	point to the parent directory	
cp	copy files <b>cp [options] &lt;source&gt; &lt;destination&gt;</b>	-r (recursive), -v (verbose)
mv	move files (works with directories too) <b>mv &lt;source&gt; &lt;destination&gt;</b>	
rm	remove files	-rf (recursive+force) (!), -i (ask before)
less	show contents of a file	
man / info	show manual about a command	
>	transfer the output from screen to a file	

# Wildcard characters

Wikipedia:

“ A wildcard character is a single character used to represent a number of characters.”

# Wildcard characters

Wikipedia:

“ A wildcard character is a single character used to represent a number of characters.”

- \* Matches zero or more characters
- ? Matches one single character
- [ ] Matches a range or a selection



# Wildcard characters

Wikipedia:

“A wildcard character is a single character used to represent a number of characters.”

\* Matches zero or more characters

? Matches one single character

[ ] Matches a range or a selection

```
[cabezon@maia run]$ ls
l                               dumpfile005000             hostfilerun~
a123                           dumpfile005500             init_scenario.f90
b111                           dumpfile006000             init_scenario.f90~
c134                           dumpfile006500             j2Msniaflame
conserveLawsRel2MpM1_08HeRot.d  dumpfile007000             nohup.out
conserveLawsRel2MpM1_08HeRot.d~ dumpfile007500             parameters.f90
createfileplot7.f90           dumpfile008000             parameters.f90~
d144                           dumpfile008500             parameters_idsa.f90
data                           dumpfile009000             readdata.f90
dumpfile000500                 dumpfile009500             readdata.f90~
dumpfile001000                 dumpfile010000             REPORT
dumpfile001500                 dumpfile010500             save
dumpfile002000                 e165                       temp
dumpfile002500                 escrituramod.f90           tiempos.d
dumpfile003000                 estabilRel2MpM1_08HeRot.d  timectrl.d
dumpfile003500                 estabilRel2MpM1_08HeRot.d~ timing.d
dumpfile004000                 find56Ni.f90
dumpfile004500                 hostfilerun
[cabezon@maia run]$
```

- Copy all **dumpfiles** into **save/**:
- Copy all **.d** files into **save/**:
- Copy all **dumpfiles** greater than **9500** into **save/**:
- Remove all **dumpfiles** between **5000** and **9500** (both included):
- Remove the five files **a123**, **b111**, **c134**, **d144**, **e165**:

# Wildcard characters

## Wikipedia:

“A wildcard character is a single character used to represent a number of characters.”

\* Matches zero or more characters

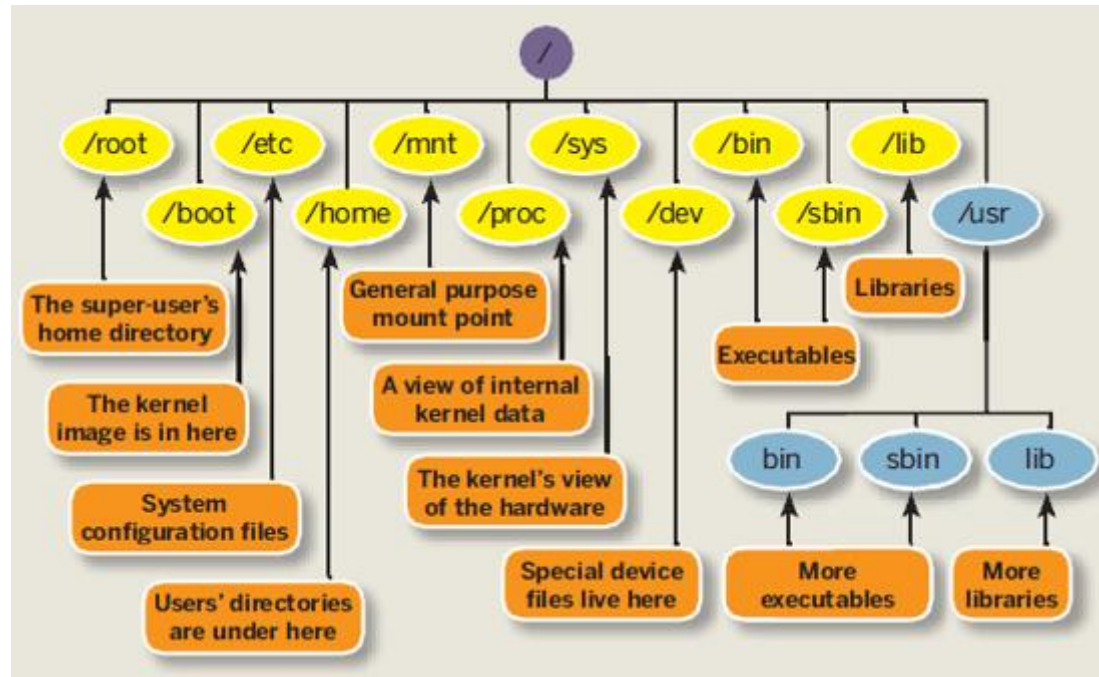
? Matches one single character

[ ] Matches a range or a selection

```
[cabezon@maia run]$ ls
l                               dumpfile005000             hostfilerun~
a123                           dumpfile005500             init_scenario.f90
b111                           dumpfile006000             init_scenario.f90~
c134                           dumpfile006500             j2Msniaflame
conserveLawsRel2MpM1_08HeRot.d dumpfile007000             nohup.out
conserveLawsRel2MpM1_08HeRot.d~ dumpfile007500             parameters.f90
createfileplot7.f90           dumpfile008000             parameters.f90~
d144                           dumpfile008500             parameters_idsa.f90
data                           dumpfile009000             readdata.f90
dumpfile000500                 dumpfile009500             readdata.f90~
dumpfile001000                 dumpfile010000             REPORT
dumpfile001500                 dumpfile010500             save
dumpfile002000                 e165                       temp
dumpfile002500                 escrituramod.f90           tiempos.d
dumpfile003000                 estabilRel2MpM1_08HeRot.d  timectrl.d
dumpfile003500                 estabilRel2MpM1_08HeRot.d~ timing.d
dumpfile004000                 find56Ni.f90
dumpfile004500                 hostfilerun
[cabezon@maia run]$
```

- Copy all **dumpfiles** into **save/**: REDACTED
- Copy all **.d** files into **save/**: REDACTED
- Copy all **dumpfiles** greater than **9500** into **save/**: REDACTED
- Remove all **dumpfiles** between **5000** and **9500** (both included): REDACTED
- Remove the five files **a123**, **b111**, **c134**, **d144**, **e165**: REDACTED

# Linux filesystem



**/home/<username>** is the place where your stuff is stored. Usual place to work in the local computer.

In the test cluster is **/shared/home/<username>**

## More about files

```
cabezon@login11:~  
-rw-r--r-- 1 cabezon physik 121 Jun 22 2012 .kshrc  
-rw----- 1 cabezon physik 234 Aug 19 16:26 .lessht  
drwxrwxr-x 3 cabezon physik 19 Jan 16 2015 .lmod.d/  
drwx----- 3 cabezon physik 18 Dec 11 2013 .local/  
drwxr-xr-x 7 cabezon physik 74 Jan 22 2015 martin/  
drwxr-xr-x 3 cabezon physik 19 Jan 24 2014 .matlab/  
drwxr-xr-x 2 cabezon physik 27 May 2 2013 .matplotlib/  
drwx----- 2 cabezon physik 58 Jan 19 2015 .mc/  
drwxr-xr-x 5 cabezon physik 51 May 2 2013 .mozilla/  
-rw-r--r-- 1 cabezon physik 2458 Sep 12 2014 new_absor_effect.f90  
-rw-r--r-- 1 cabezon physik 18240 May 8 11:55 nuprox3D.f90  
drwxr----- 3 cabezon physik 18 May 2 2013 .pki/  
-rw-r--r-- 1 cabezon physik 6525 Jun 23 12:05 progenitor.f  
-rw----- 1 cabezon physik 256 May 2 2013 .pulse-cookie  
-rw-r--r-- 1 cabezon physik 1092 Jan 22 2015 REPORT  
drwxr-xr-x 21 cabezon physik 4096 Aug 25 09:33 research/  
lrwxrwxrwx 1 root root 13 Nov 16 2012 save -> /save/cabezon  
drwxr-xr-x 3 cabezon physik 19 Jan 12 2015 software/  
drwx----- 2 cabezon physik 95 Jan 16 2015 .ssh/  
drwxr-xr-x 6 cabezon physik 4096 Jan 22 2015 tmp/  
-rw----- 1 cabezon physik 8260 Aug 20 17:18 .viminfo  
-rw----- 1 cabezon physik 13390 Aug 26 14:53 .Xauthority  
-rw-r--r-- 1 cabezon physik 658 Jun 23 2012 .zshrc  
cabezon@login11 ~]$
```

Type: **d** (directory), **-** (ordinary file), **l** (symbolic link)

## More about files

```
cabezon@login11:~  
-rw-r--r-- 1 cabezon physik 121 Jun 22 2012 .kshrc  
-rw----- 1 cabezon physik 234 Aug 19 16:26 .lessht  
drwxrwxr-x 3 cabezon physik 19 Jan 16 2015 .lmod.d/  
drwx----- 3 cabezon physik 18 Dec 11 2013 .local/  
drwxr-xr-x 7 cabezon physik 74 Jan 22 2015 martin/  
drwxr-xr-x 3 cabezon physik 19 Jan 24 2014 .matlab/  
drwxr-xr-x 2 cabezon physik 27 May 2 2013 .matplotlib/  
drwx----- 2 cabezon physik 58 Jan 19 2015 .mc/  
drwxr-xr-x 5 cabezon physik 51 May 2 2013 .mozilla/  
-rw-r--r-- 1 cabezon physik 2458 Sep 12 2014 new_absor_effect.f90  
-rw-r--r-- 1 cabezon physik 18240 May 8 11:55 nuprox3D.f90  
drwxr----- 3 cabezon physik 18 May 2 2013 .pki/  
-rw-r--r-- 1 cabezon physik 6525 Jun 23 12:05 progenitor.f  
-rw----- 1 cabezon physik 256 May 2 2013 .pulse-cookie  
-rw-r--r-- 1 cabezon physik 1092 Jan 22 2015 REPORT  
drwxr-xr-x 21 cabezon physik 4096 Aug 25 09:33 research/  
lrwxrwxrwx 1 root root 13 Nov 16 2012 save -> /save/cabezon  
drwxr-xr-x 3 cabezon physik 19 Jan 12 2015 software/  
drwx----- 2 cabezon physik 95 Jan 16 2015 .ssh/  
drwxr-xr-x 6 cabezon physik 4096 Jan 22 2015 tmp/  
-rw----- 1 cabezon physik 8260 Aug 20 17:18 .viminfo  
-rw----- 1 cabezon physik 13390 Aug 26 14:53 .Xauthority  
-rw-r--r-- 1 cabezon physik 658 Jun 23 2012 .zshrc  
[cabezon@login11 ~]$
```

Permissions: define the access rights.

## More about files

```
cabezon@login11:~  
-rw-r--r-- 1 cabezon physik 121 Jun 22 2012 .kshrc  
-rw----- 1 cabezon physik 234 Aug 19 16:26 .lessht  
drwxrwxr-x 3 cabezon physik 19 Jan 16 2015 .lmod.d/  
drwx----- 3 cabezon physik 18 Dec 11 2013 .local/  
drwxr-xr-x 7 cabezon physik 74 Jan 22 2015 martin/  
drwxr-xr-x 3 cabezon physik 19 Jan 24 2014 .matlab/  
drwxr-xr-x 2 cabezon physik 27 May 2 2013 .matplotlib/  
drwx----- 2 cabezon physik 58 Jan 19 2015 .mc/  
drwxr-xr-x 5 cabezon physik 51 May 2 2013 .mozilla/  
-rw-r--r-- 1 cabezon physik 2458 Sep 12 2014 new_absor_effect.f90  
-rw-r--r-- 1 cabezon physik 18240 May 8 11:55 nuprox3D.f90  
drwxr----- 3 cabezon physik 18 May 2 2013 .pki/  
-rw-r--r-- 1 cabezon physik 6525 Jun 23 12:05 progenitor.f  
-rw----- 1 cabezon physik 256 May 2 2013 .pulse-cookie  
-rw-r--r-- 1 cabezon physik 1092 Jan 22 2015 REPORT  
drwxr-xr-x 21 cabezon physik 4096 Aug 25 09:33 research/  
lrwxrwxrwx 1 root root 13 Nov 16 2012 save -> /save/cabezon  
drwxr-xr-x 3 cabezon physik 19 Jan 12 2015 software/  
drwx----- 2 cabezon physik 95 Jan 16 2015 .ssh/  
drwxr-xr-x 6 cabezon physik 4096 Jan 22 2015 tmp/  
-rw----- 1 cabezon physik 8260 Aug 20 17:18 .viminfo  
-rw----- 1 cabezon physik 13390 Aug 26 14:53 .Xauthority  
-rw-r--r-- 1 cabezon physik 658 Jun 23 2012 .zshrc  
[cabezon@login11 ~]$
```

Links: number of links pointing to this file/directory.

## More about files

```
cabezon@login11:~  
-rw-r--r-- 1 cabezon physik 121 Jun 22 2012 .kshrc  
-rw----- 1 cabezon physik 234 Aug 19 16:26 .lessht  
drwxrwxr-x 3 cabezon physik 19 Jan 16 2015 .lmod.d/  
drwx----- 3 cabezon physik 18 Dec 11 2013 .local/  
drwxr-xr-x 7 cabezon physik 74 Jan 22 2015 martin/  
drwxr-xr-x 3 cabezon physik 19 Jan 24 2014 .matlab/  
drwxr-xr-x 2 cabezon physik 27 May 2 2013 .matplotlib/  
drwx----- 2 cabezon physik 58 Jan 19 2015 .mc/  
drwxr-xr-x 5 cabezon physik 51 May 2 2013 .mozilla/  
-rw-r--r-- 1 cabezon physik 2458 Sep 12 2014 new_absor_effect.f90  
-rw-r--r-- 1 cabezon physik 18240 May 8 11:55 nuprox3D.f90  
drwxr----- 3 cabezon physik 18 May 2 2013 .pki/  
-rw-r--r-- 1 cabezon physik 6525 Jun 23 12:05 progenitor.f  
-rw----- 1 cabezon physik 256 May 2 2013 .pulse-cookie  
-rw-r--r-- 1 cabezon physik 1092 Jan 22 2015 REPORT  
drwxr-xr-x 21 cabezon physik 4096 Aug 25 09:33 research/  
lrwxrwxrwx 1 root root 13 Nov 16 2012 save -> /save/cabezon  
drwxr-xr-x 3 cabezon physik 19 Jan 12 2015 software/  
drwx----- 2 cabezon physik 95 Jan 16 2015 .ssh/  
drwxr-xr-x 6 cabezon physik 4096 Jan 22 2015 tmp/  
-rw----- 1 cabezon physik 8260 Aug 20 17:18 .viminfo  
-rw----- 1 cabezon physik 13390 Aug 26 14:53 .Xauthority  
-rw-r--r-- 1 cabezon physik 658 Jun 23 2012 .zshrc  
[cabezon@login11 ~]$
```

→ Owner: Usually the user who created the file.



## More about files

```
cabezon@login11:~  
-rw-r--r-- 1 cabezon physik 121 Jun 22 2012 .kshrc  
-rw----- 1 cabezon physik 234 Aug 19 16:26 .lessht  
drwxrwxr-x 3 cabezon physik 19 Jan 16 2015 .lmod.d/  
drwx----- 3 cabezon physik 18 Dec 11 2013 .local/  
drwxr-xr-x 7 cabezon physik 74 Jan 22 2015 martin/  
drwxr-xr-x 3 cabezon physik 19 Jan 24 2014 .matlab/  
drwxr-xr-x 2 cabezon physik 27 May 2 2013 .matplotlib/  
drwx----- 2 cabezon physik 58 Jan 19 2015 .mc/  
drwxr-xr-x 5 cabezon physik 51 May 2 2013 .mozilla/  
-rw-r--r-- 1 cabezon physik 2458 Sep 12 2014 new_absor_effect.f90  
-rw-r--r-- 1 cabezon physik 18240 May 8 11:55 nuprox3D.f90  
drwxr----- 3 cabezon physik 18 May 2 2013 .pki/  
-rw-r--r-- 1 cabezon physik 6525 Jun 23 12:05 progenitor.f  
-rw----- 1 cabezon physik 256 May 2 2013 .pulse-cookie  
-rw-r--r-- 1 cabezon physik 1092 Jan 22 2015 REPORT  
drwxr-xr-x 21 cabezon physik 4096 Aug 25 09:33 research/  
lrwxrwxrwx 1 root root 13 Nov 16 2012 save -> /save/cabezon  
drwxr-xr-x 3 cabezon physik 19 Jan 12 2015 software/  
drwx----- 2 cabezon physik 95 Jan 16 2015 .ssh/  
drwxr-xr-x 6 cabezon physik 4096 Jan 22 2015 tmp/  
-rw----- 1 cabezon physik 8260 Aug 20 17:18 .viminfo  
-rw----- 1 cabezon physik 13390 Aug 26 14:53 .Xauthority  
-rw-r--r-- 1 cabezon physik 658 Jun 23 2012 .zshrc  
[cabezon@login11 ~]$
```

Group: Set of users that can access the file according to the permissions specified in the group field.



## More about files

```
cabezon@login11:~  
-rw-r--r-- 1 cabezon physik 121 Jun 22 2012 .kshrc  
-rw----- 1 cabezon physik 234 Aug 19 16:26 .lessht  
drwxrwxr-x 3 cabezon physik 19 Jan 16 2015 .lmod.d/  
drwx----- 3 cabezon physik 18 Dec 11 2013 .local/  
drwxr-xr-x 7 cabezon physik 74 Jan 22 2015 martin/  
drwxr-xr-x 3 cabezon physik 19 Jan 24 2014 .matlab/  
drwxr-xr-x 2 cabezon physik 27 May 2 2013 .matplotlib/  
drwx----- 2 cabezon physik 58 Jan 19 2015 .mc/  
drwxr-xr-x 5 cabezon physik 51 May 2 2013 .mozilla/  
-rw-r--r-- 1 cabezon physik 2458 Sep 12 2014 new_absor_effect.f90  
-rw-r--r-- 1 cabezon physik 18240 May 8 11:55 nuprox3D.f90  
drwxr----- 3 cabezon physik 18 May 2 2013 .pki/  
-rw-r--r-- 1 cabezon physik 6525 Jun 23 12:05 progenitor.f  
-rw----- 1 cabezon physik 256 May 2 2013 .pulse-cookie  
-rw-r--r-- 1 cabezon physik 1092 Jan 22 2015 REPORT  
drwxr-xr-x 21 cabezon physik 4096 Aug 25 09:33 research/  
lrwxrwxrwx 1 root root 13 Nov 16 2012 save -> /save/cabezon  
drwxr-xr-x 3 cabezon physik 19 Jan 12 2015 software/  
drwx----- 2 cabezon physik 95 Jan 16 2015 .ssh/  
drwxr-xr-x 6 cabezon physik 4096 Jan 22 2015 tmp/  
-rw----- 1 cabezon physik 8260 Aug 20 17:18 .viminfo  
-rw----- 1 cabezon physik 13390 Aug 26 14:53 .Xauthority  
-rw-r--r-- 1 cabezon physik 658 Jun 23 2012 .zshrc  
[cabezon@login11 ~]$
```

→ Size: expressed in bytes. ( **-h** option is more ergonomic)

## More about files

```
cabezon@login11:~  
-rw-r--r-- 1 cabezon physik 121 Jun 22 2012 .kshrc  
-rw----- 1 cabezon physik 234 Aug 19 16:26 .lessht  
drwxrwxr-x 3 cabezon physik 19 Jan 16 2015 .lmod.d/  
drwx----- 3 cabezon physik 18 Dec 11 2013 .local/  
drwxr-xr-x 7 cabezon physik 74 Jan 22 2015 martin/  
drwxr-xr-x 3 cabezon physik 19 Jan 24 2014 .matlab/  
drwxr-xr-x 2 cabezon physik 27 May 2 2013 .matplotlib/  
drwx----- 2 cabezon physik 58 Jan 19 2015 .mc/  
drwxr-xr-x 5 cabezon physik 51 May 2 2013 .mozilla/  
-rw-r--r-- 1 cabezon physik 2458 Sep 12 2014 new_absor_effect.f90  
-rw-r--r-- 1 cabezon physik 18240 May 8 11:55 nuprox3D.f90  
drwxr----- 3 cabezon physik 18 May 2 2013 .pki/  
-rw-r--r-- 1 cabezon physik 6525 Jun 23 12:05 progenitor.f  
-rw----- 1 cabezon physik 256 May 2 2013 .pulse-cookie  
-rw-r--r-- 1 cabezon physik 1092 Jan 22 2015 REPORT  
drwxr-xr-x 21 cabezon physik 4096 Aug 25 09:33 research/  
lrwxrwxrwx 1 root root 13 Nov 16 2012 save -> /save/cabezon  
drwxr-xr-x 3 cabezon physik 19 Jan 12 2015 software/  
drwx----- 2 cabezon physik 95 Jan 16 2015 .ssh/  
drwxr-xr-x 6 cabezon physik 4096 Jan 22 2015 tmp/  
-rw----- 1 cabezon physik 8260 Aug 20 17:18 .viminfo  
-rw----- 1 cabezon physik 13390 Aug 26 14:53 .Xauthority  
-rw-r--r-- 1 cabezon physik 658 Jun 23 2012 .zshrc  
[cabezon@login11 ~]$
```

Date: Date and time of last modification (write).

The **-u** option displays time of last access (read).

## More about files

```
cabezon@login11:~  
-rw-r--r-- 1 cabezon physik 121 Jun 22 2012 .kshrc  
-rw----- 1 cabezon physik 234 Aug 19 16:26 .lessht  
drwxrwxr-x 3 cabezon physik 19 Jan 16 2015 .lmod.d/  
drwx----- 3 cabezon physik 18 Dec 11 2013 .local/  
drwxr-xr-x 7 cabezon physik 74 Jan 22 2015 martin/  
drwxr-xr-x 3 cabezon physik 19 Jan 24 2014 .matlab/  
drwxr-xr-x 2 cabezon physik 27 May 2 2013 .matplotlib/  
drwx----- 2 cabezon physik 58 Jan 19 2015 .mc/  
drwxr-xr-x 5 cabezon physik 51 May 2 2013 .mozilla/  
-rw-r--r-- 1 cabezon physik 2458 Sep 12 2014 new_absor_effect.f90  
-rw-r--r-- 1 cabezon physik 18240 May 8 11:55 nuprox3D.f90  
drwxr----- 3 cabezon physik 18 May 2 2013 .pki/  
-rw-r--r-- 1 cabezon physik 6525 Jun 23 12:05 progenitor.f  
-rw----- 1 cabezon physik 256 May 2 2013 .pulse-cookie  
-rw-r--r-- 1 cabezon physik 1092 Jan 22 2015 REPORT  
drwxr-xr-x 21 cabezon physik 4096 Aug 25 09:33 research/  
lrwxrwxrwx 1 root root 13 Nov 16 2012 save -> /save/cabezon  
drwxr-xr-x 3 cabezon physik 19 Jan 12 2015 software/  
drwx----- 2 cabezon physik 95 Jan 16 2015 .ssh/  
drwxr-xr-x 6 cabezon physik 4096 Jan 22 2015 tmp/  
-rw----- 1 cabezon physik 8260 Aug 20 17:18 .viminfo  
-rw----- 1 cabezon physik 13390 Aug 26 14:53 .Xauthority  
-rw-r--r-- 1 cabezon physik 658 Jun 23 2012 .zshrc  
[cabezon@login11 ~]$
```

→ Name: this is self-explaining...

# More about rights

cabezon@login11:~

There are 9 characters grouped in 3 sets of 3.  
Each set defines permissions to 3 access types: **r** (read), **w** (write), **x** (execute)  
An hyphen **-** means no permission to that access type.

```
rw-r--r-- 1 cabezon physik 121 Jun 22 2012 .kshrc
rw----- 1 cabezon physik 234 Jan 19 15:26 .lessht
rwxrwxr-x 1 cabezon physik 19 Jan 16 2015 .lmod.d/
rwx----- 1 cabezon physik 19 Jan 16 2015 .lmod.d/
rwxr-xr-x 1 cabezon physik 19 Jan 16 2015 .lmod.d/
rwxr-xr-x 3 cabezon physik 19 Jan 24 2014 .matlab/
rwxr-xr-x 2 cabezon physik 27 May 2 2013 .matplotlib/
rwx----- 2 cabezon physik 58 Jan 19 2015 .mc/
rwxr-xr-x 5 cabezon physik 51 May 2 2013 .mozilla/
rw-r--r-- 1 cabezon physik 2458 Sep 12 2014 new_absor_effect.f90
rw-r--r-- 1 cabezon physik 18240 May 8 11:55 nuprox3D.f90
rwxr----- 3 cabezon physik 18 May 2 2013 .pki/
rw-r--r-- 1 cabezon physik 6525 Jun 23 12:05 progenitor.f
rw----- 1 cabezon physik 256 May 2 2013 .pulse-cookie
rw-r--r-- 1 cabezon physik 1092 Jan 22 2015 REPORT
rwxr-xr-x 21 cabezon physik 4096 Aug 25 09:33 research/
rwxrwxrwx 1 root root 13 Nov 16 2012 save -> /save/cabezon
rwxr-xr-x 3 cabezon physik 19 Jan 12 2015 software/
rwx----- 2 cabezon physik 95 Jan 16 2015 .ssh/
rwxr-xr-x 6 cabezon physik 4096 Jan 22 2015 tmp/
rw----- 1 cabezon physik 8260 Aug 20 17:18 .viminfo
rw----- 1 cabezon physik 13390 Aug 26 14:53 .Xauthority
rw-r--r-- 1 cabezon physik 658 Jun 23 2012 .zshrc
```

[cabezon@login11 ~]\$

# More about rights

```
cabezon@login11:~  
rw-r--r-- 1 cabezon physik 121 Jun 22 2012 .kshrc  
rw-r--r-- 1 cabezon physik 234 Jan 19 2015 .lessht  
rwxrwxr-x 1 cabezon physik 19 Jan 16 2015 .lmod.d/  
rwx----- 1 cabezon physik 19 Jan 16 2015 .lmod.rc  
rwxr-xr-x 3 cabezon physik 19 Jan 24 2014 .matlab/  
rwxr-xr-x 1 cabezon physik 19 Jan 16 2015 .matplotlib/  
rwxr-xr-x 2 cabezon physik 58 Jan 19 2015 .mo/  
rwxr-xr-x 5 cabezon physik 51 May 2 2013 .mozilla/  
rw-r--r-- 1 cabezon physik 2458 Sep 12 2014 new_absor_effect.f90  
rw-r--r-- 1 cabezon physik 18240 May 8 11:55 nuprox3D.f90  
rwxr----- 3 cabezon physik 18 May 2 2013 .pki/  
rw-r--r-- 1 cabezon physik 6525 Jun 23 12:05 progenitor.f  
rw----- 1 cabezon physik 256 May 2 2013 .pulse-cookie  
rw-r--r-- 1 cabezon physik 1092 Jan 22 2015 REPORT  
rwxr-xr-x 21 cabezon physik 4096 Aug 25 09:33 research/  
rwxrwxrwx 1 root root 13 Nov 16 2012 save -> /save/cabezon  
rwxr-xr-x 3 cabezon physik 19 Jan 12 2015 software/  
rwx----- 2 cabezon physik 95 Jan 16 2015 .ssh/  
rwxr-xr-x 6 cabezon physik 4096 Jan 22 2015 tmp/  
rw----- 1 cabezon physik 8260 Aug 20 17:18 .viminfo  
rw----- 1 cabezon physik 13390 Aug 26 14:53 .Xauthority  
rw-r--r-- 1 cabezon physik 658 Jun 23 2012 .zshrc  
cabezon@login11 ~]$
```

There are 9 characters grouped in 3 sets of 3.

Each set defines permissions to 3 access types: **r** (read), **w** (write), **x** (execute)

An hyphen **-** means no permission to that access type.

Each set refers to different users: **owner**, **group**, **other**

# More about rights

cabezon@login11:~

```
rw-r--r--
rw-----
rwxrwxr-x
rwx-----
rwxr-xr-x
rwxr-xr-x
rwxr-xr-x
rwx-----
rwxr-xr-x
rw-r--r--
rw-r--r--
rw-r--r--
rwxr-xr-x
rwxrwxrwx
rwxr-xr-x
rwx-----
rwxr-xr-x
rw-----
rw-----
rw-r--r--
```

There are 9 characters grouped in 3 sets of 3.

Each set defines permissions to 3 access types: **r** (read), **w** (write), **x** (execute)

An hyphen **-** means no permission to that access type.

Each set refers to different users: **owner**, **group**, **other**

Example:

```
1 cabezon fisik 121 Jun 22 2012 .kshrc
1 cabezon fisik 234 Jan 19 2013 .lessht
1 cabezon fisik 19 Jan 16 2015 .lmod.d/
3 cabezon fisik 19 Jan 24 2014 .matlab/
1 cabezon fisik 19 Jan 19 2015 .matplotlib/
2 cabezon fisik 58 Jan 19 2015 .mo/
1 cabezon fisik 51 May 2 2013 .mozilla/
1 cabezon fisik 2458 Sep 12 2014 new_absor_effect.f90
1 cabezon fisik 18240 May 8 11:55 nuprox3D.f90
3 cabezon fisik 18 May 2 2013 .pki/
1 cabezon fisik 6525 Jun 23 12:05 progenitor.f
1 cabezon fisik 256 May 2 2013 .pulse-cookie
1 cabezon fisik 1092 Jan 22 2015 REPORT
21 cabezon fisik 4096 Aug 25 09:33 research/
1 root root 13 Nov 16 2012 save -> /save/cabezon
3 cabezon fisik 19 Jan 12 2015 software/
2 cabezon fisik 95 Jan 16 2015 .ssh/
6 cabezon fisik 4096 Jan 22 2015 tmp/
1 cabezon fisik 8260 Aug 20 17:18 .viminfo
1 cabezon fisik 13390 Aug 26 14:53 .Xauthority
1 cabezon fisik 658 Jun 23 2012 .zshrc
[cabezon@login11 ~]$
```

# More about rights

cabezon@login11:~

```
rw-r--r--
rw-----
rwxrwxr-x
rwx-----
rwxr-xr-x
rwxr-xr-x
rwxr-xr-x
rwx-----
rwxr-xr-x
rw-r--r--
rw-r--r--
rwxr-----
rw-r--r--
rw-----
rw-r--r--
rwxr-xr-x
rwxrwxrwx
rwxr-xr-x
rwx-----
rwxr-xr-x
rw-----
rw-----
rw-r--r--
```

There are 9 characters grouped in 3 sets of 3.

Each set defines permissions to 3 access types: **r** (read), **w** (write), **x** (execute)

An hyphen **-** means no permission to that access type.

Each set refers to different users: **owner**, **group**, **other**

Example: REDACTED

```
1 cabezon fisik 121 Jun 22 2012 .kshrc
1 cabezon fisik 334 Jan 19 2013 .lessht
1 cabezon fisik 19 Jan 12 2015 .lmod.d/
1 cabezon fisik 19 Jan 12 2015 .in/
3 cabezon fisik 19 Jan 24 2014 .matlab/
1 cabezon fisik 19 Jan 19 2015 .matplotlib/
2 cabezon fisik 58 Jan 19 2015 .mo/
1 cabezon fisik 51 May 2 2013 .mozilla/
1 cabezon fisik 2458 Sep 12 2014 new_absor_effect.f90
1 cabezon fisik 18240 May 8 11:55 nuprox3D.f90
3 cabezon fisik 18 May 2 2013 .pki/
1 cabezon fisik 6525 Jun 23 12:05 progenitor.f
1 cabezon fisik 256 May 2 2013 .pulse-cookie
1 cabezon fisik 1092 Jan 22 2015 REPORT
21 cabezon fisik 4096 Aug 25 09:33 research/
1 root root 13 Nov 16 2012 save -> /save/cabezon
3 cabezon fisik 19 Jan 12 2015 software/
2 cabezon fisik 95 Jan 16 2015 .ssh/
6 cabezon fisik 4096 Jan 22 2015 tmp/
1 cabezon fisik 8260 Aug 20 17:18 .viminfo
1 cabezon fisik 13390 Aug 26 14:53 .Xauthority
1 cabezon fisik 658 Jun 23 2012 .zshrc
```

[cabezon@login11 ~]\$

# More about rights

cabezon@login11:~

```
rw-r--r--
rw-----
rwxrwxr-x
rwx-----
rwxr-xr-x
rwxr-xr-x
rwxr-xr-x
rwx-----
rwxr-xr-x
rw-r--r--
rw-r--r--
rwxr-----
rw-r--r--
rw-----
rw-r--r--
rwxr-xr-x
rwxrwxrwx
rwxr-xr-x
rwx-----
rwxr-xr-x
rw-----
rw-----
rw-r--r--
```

There are 9 characters grouped in 3 sets of 3.

Each set defines permissions to 3 access types: **r** (read), **w** (write), **x** (execute)

An hyphen **-** means no permission to that access type.

Each set refers to different users: **owner**, **group**, **other**

Example: REDACTED

**chmod**: used to change permissions.

**Symbolic**: uses **u**, **g**, **o**, **a** an operator **+**, **-** and a mode **r**, **w**, **x**

Ex: **chmod g+w <file>** : gives write permission to the group users

**chmod -R a-wx <directory>** : removes writing and executing rights to all users, recursively for the directory tree inside **<directory>**

**Octal**: uses three digits to define a combination of rights for each set. It's simply the decimal number corresponding to the binary resulting from the set of permissions.

Ex: **--x = 001 = 1**, **r-x = 101 = 5**, **rwx = 111 = 7**

**chmod 662 <file> =**



# More about rights

cabezon@login11:~

```
rw-r--r--
rw-----
rwxrwxr-x
rwx-----
rwxr-xr-x
rwxr-xr-x
rwxr-xr-x
rwx-----
rwxr-xr-x
rw-r--r--
rw-r--r--
rwxr-----
rw-r--r--
rw-----
rw-r--r--
rwxr-xr-x
rwxrwxrwx
rwxr-xr-x
rwx-----
rwxr-xr-x
rw-----
rw-----
rw-r--r--
```

There are 9 characters grouped in 3 sets of 3.

Each set defines permissions to 3 access types: **r** (read), **w** (write), **x** (execute)

An hyphen **-** means no permission to that access type.

Each set refers to different users: **owner**, **group**, **other**

Example: REDACTED

**chmod**: used to change permissions.

**Symbolic**: uses **u**, **g**, **o**, **a** an operator **+**, **-** and a mode **r**, **w**, **x**

Ex: **chmod g+w <file>** : gives write permission to the group users

**chmod -R a-wx <directory>** : removes writing and executing rights to all users, recursively for the directory tree inside **<directory>**

**Octal**: uses three digits to define a combination of rights for each set. It's simply the decimal number corresponding to the binary resulting from the set of permissions.

Ex: **--x = 001 = 1**, **r-x = 101 = 5**, **rwx = 111 = 7**

**chmod 662 <file> = REDACTED**

There are only 10 types of people in this world:  
those who understand binary and those who don't.

# More about ownership

cabezon@login11:~

```
-rw-r--r-- 1 cabezon physik
-rw----- 1 cabezon physik
drwxrwxr-x 3 cabezon physik
drwx----- 3 cabezon physik
drwxr-xr-x 7 cabezon physik
drwxr-xr-x 3 cabezon physik
drwxr-xr-x 2 cabezon physik
drwx----- 2 cabezon physik
drwxr-xr-x 5 cabezon physik
-rw-r--r-- 1 cabezon physik
-rw-r--r-- 1 cabezon physik
drwxr----- 3 cabezon physik
-rw-r--r-- 1 cabezon physik
-rw----- 1 cabezon physik
-rw-r--r-- 1 cabezon physik
drwxr-xr-x 21 cabezon physik
lrwxrwxrwx 1 root root
drwxr-xr-x 3 cabezon physik
drwx----- 2 cabezon physik
drwxr-xr-x 6 cabezon physik
-rw----- 1 cabezon physik
-rw----- 1 cabezon physik
-rw-r--r-- 1 cabezon physik
[cabezon@login11 ~]$
```

Usually our home is blocked to other users, unless we have allowed read permission to users that are share a group with us.

This is a way of safely sharing data by allowing access.

**chown:** allows to change the ownership of a file/directory

**Ex: chown alice <file>** : changes ownership to **alice**

**chown -R alice:physik <directory>** : changes recursively ownership to **alice** and group to **physik**.

```
121 Jun 22 2012 .kshrc
334 Aug 19 15:25 .ssh/
10 Jan 16 2015 .viminfo
74 Jan 22 2015 martin/
19 Jan 24 2014 .matlab/
58 Jan 19 2015 .mo/
51 May 2 2013 .mozilla/
18240 May 8 11:55 nuprox3D.f90
4096 Aug 25 09:33 research/
13 Nov 16 2012 save -> /save/cabezon
19 Jan 12 2015 software/
95 Jan 16 2015 .ssh/
4096 Jan 22 2015 tmp/
8260 Aug 20 17:18 .viminfo
13390 Aug 26 14:53 .Xauthority
658 Jun 23 2012 .zshrc
```

## 2. Exercise

1. Log in to the cluster
2. Check your path: `pwd`. That's your home directory.
3. Check the content of the environment variable `$HOME`: `echo $HOME`
4. Enter in `linux_course/` and create a directory named `Exercise2/`
5. Create a file: `cat > hello1.txt`
6. Write `hello world` and exit (`CTRL+D` twice)
7. Create another file. Do `emacs hello2.txt`
8. Write something inside, save it (`CTRL+X`, `CTRL+S`) and exit (`CTRL+X`, `CTRL+C`)
9. Create another file: `nano hello3.txt`
10. Write `hello world again` and exit (`CTRL+X`). Save when asked.
11. Create another file: `vim hello4.txt`
12. Press `i` to go to insert mode. Write something. Press `ESC` and exit (`:wq`)
13. Write `cp he` and press `TAB`. What happens? Press `TAB` twice. Add a `4` and press `TAB`. What happens?
14. Finish the command to copy `hello4.txt` into `hello5.txt`
15. Go to your `home/` with `cd`
16. Find where all `hello` files are: `find linux_course/ -name 'hello*'`
17. From your home directory and with one command move all `hello` files from `linux_course/` to `linux_course/Exercise2/`
18. Go to `Exercise2/` and check that all 5 `hello` files are there.

The final outcome should be a directory structure `linux_course/Exercise2` that includes 5 files named `hello1.txt`, `hello2.txt`,...

## 2. Exercise

### Summary

Command	Used to...	Common options
pwd	print current working directory	
echo	output text to screen	
cat	displays/create file contents	
emacs	displays/create file contents	-nw (no window)
nano	displays/create file contents	
vim	displays/create file contents	
find	search for a file	<path> -name '<file>'
<TAB>	auto completion	
\$variable	dynamic values that can control how processes behave	\$HOME, \$PATH, \$TMPDIR, \$LD_LIBRARY_PATH, \$PWD

## 2. Exercise

### VIM

usable in just about  
any environment.

does one thing, well.



Used to...	Common options
Change current working directory	
Print text to screen	
Displays/create file contents	
Displays/create file contents	-nw (no window)
Displays/create file contents	
Displays/create file contents	
Search for a file	<path> -name '<file>'
Tab completion	
Dynamic values that can control how processes behave	\$HOME, \$PATH, \$TMPDIR, \$LD_LIBRARY_PATH, \$PWD

## 2. Exercise

### VIM

usable in just about  
any environment.

does one thing, well.



### EMACS

flexible, customizable, and  
packed with every feature  
known to man.



### Common options

(no window)

th> -name '<file>'

ME, \$PATH, \$TMPDIR,  
\_LIBRARY\_PATH, \$PWD

## 2. Exercise

### VIM

usable in just about  
any environment.

does one thing, well.



### EMACS

flexible, customizable, and  
packed with every feature  
known to man.



### NANO

mostly used by people  
who do not know  
what they are doing;  
or psychopaths.



© 2015 CURTIS LASSAM - CUBE-DRONE.COM

Want to learn vim? Check <http://vim-adventures.com>

### 3. Exercise

1. Log in to the cluster
2. Enter in `linux_course/` and create a directory named `Exercise3/`
3. Find in which files from `Exercise2/` is 'hello' written: `grep 'hello' Exercise2/*`
4. Compress the `Exercise2/` directory: `tar czvf exercise2.tar.gz Exercise2/`
5. Move `exercise2.tar.gz` to `Exercise3/`
6. Enter in `Exercise3/`
7. Extract the contents of the tar ball: `tar xvf exercise2.tar.gz`
8. Check the contents of `Exercise3/`
9. Rename the newly created folder `Exercise2/` to `data/`
10. Enter in `data/` and check the size of all files: `ll` (these are two letters!)
11. Compress all txt files : `gzip *.txt` and check the size of the files. Comments?
12. Extract all compressed files: `gunzip *.gz`
13. Get this data file:  
`cp /shared/data/linux/tau_nue70000.d .`
14. Check its contents with `less`.
15. Check size of the data file, compress it and check it again.
16. Copy and rename the `.gz` file as `data.d.gz` with one single instruction.
17. Extract `data.d.gz`. Note that the uncompressed file has also the new name.
18. Extract `tau_nue70000.d.gz`
19. Check that both files are the same: `diff -s data.d tau_nue70000.d`
20. Edit `data.d`, change one number, save it, and diff again.
21. Print the second and third columns of `data.d` in a file named `out.d`:  
`awk '{print $2 " " $3}' data.d > out.d`
22. Check the contents of `out.d` and delete `data.d`
23. Go to your home and check the size of `linux_course/`: `du -h linux_course`



## 3. Exercise

### Summary

Command	Used to...	Common options
grep	search file (or stdin) for lines matching a pattern	-i (case insensitive), -w (only)
tar	back up entire directories and file into a single container file	-c (create), -v (verbose), -f (file), -z (zip), -x (extract), -t (list)
gzip	compress individual files	-d (decompress)
gunzip	decompress individual files	
diff	display differences between two files	-s (report if files are identical)
awk	manipulate files with data in columns	
du	show the disk usage of a directory	-h (human readable) -s (display only the sum)

## Extra tips

You can use **&** at the end of a command to keep using the command line and allow the execution to proceed in the background. Ex: **emacs datafile &**

**time** in front of any command will provide the time it takes to execute. Ex: **time /bin/sleep 5**

**history** shows a numbered list of the commands that have been used. With **!<number>** that command is re-used. With CTRL+R it performs a reverse search in the history.

**top** gives you information about running processes, memory and CPU consumption in realtime.

**htop** is an enhanced version of **top**, which also gives a pseudo-graphical load / CPU.

**ps aux** lists all processes, like **top**, but as a snapshot.

**\$HOME** is an environment variable that points to your **home/** directory.

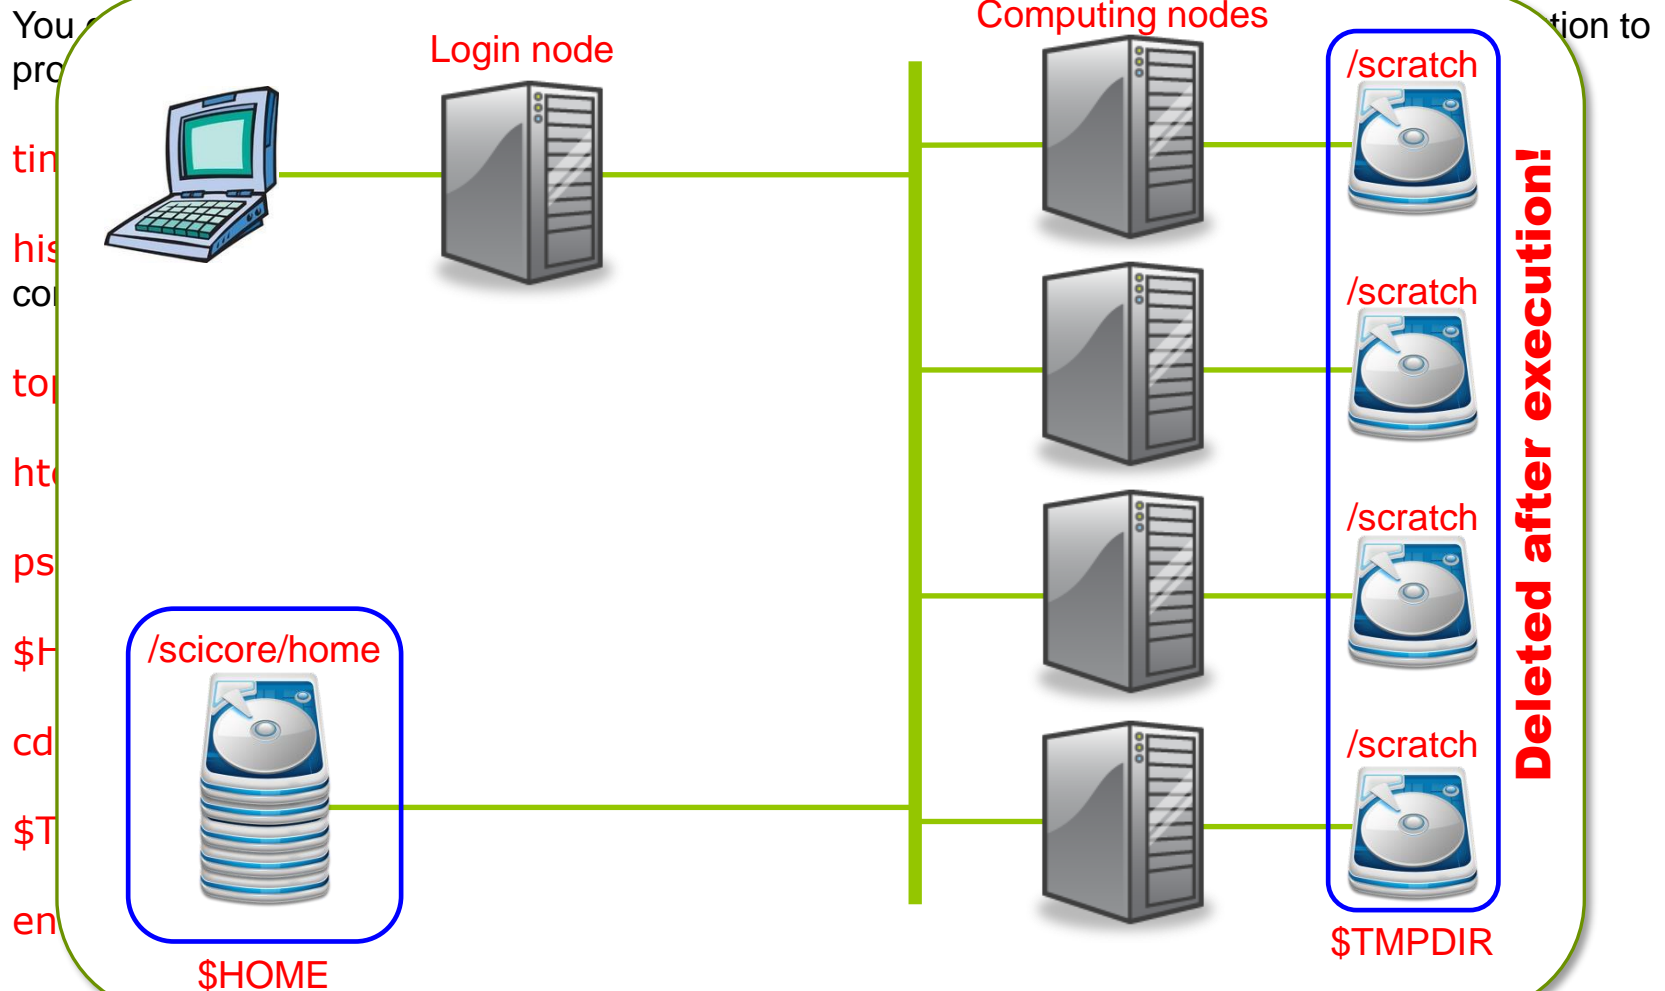
**cd** without arguments will always lead you to **cd \$HOME**

**\$TMPDIR** points to a temporary directory

**env** will show all environment variables

Copy and paste is very easy with a mouse: if you select text it is already copied. To paste it press the central button of the mouse (if you have one...) or right-click and paste.

## Extra tips



Copy and paste is very easy with a mouse. If you select text it is already copied. To paste it press the central button of the mouse (if you have one...) or right-click and paste.

## Extra tips

Event	Latency	Scaled
1 CPU cycle	0.3 ns	1 s
Level 1 cache access	0.9 ns	3 s
Level 2 cache access	2.8 ns	9 s
Level 3 cache access	12.9 ns	43 s
Main memory access (DRAM, from CPU)	120 ns	6 min
Solid-state disk I/O (flash memory)	50–150 $\mu$ s	2–6 days
Rotational disk I/O	1–10 ms	1–12 months
Internet: San Francisco to New York	40 ms	4 years
Internet: San Francisco to United Kingdom	81 ms	8 years
Internet: San Francisco to Australia	183 ms	19 years
TCP packet retransmit	1–3 s	105–317 years
OS virtualization system reboot	4 s	423 years
SCSI command time-out	30 s	3 millennia
Hardware (HW) virtualization system reboot	40 s	4 millennia
Physical system reboot	5 m	32 millennia

## Extra tips

You can use **grep** attached to any command with the pipe **|**. This will allow you to show only the information you are interested in. Ex: **history** vs **history | grep ls**

In fact, pipes can be used with almost all commands. They transfer the output of the first command as input for the second one, and so on.

Ex: **ps aux | grep -w grep | awk '{print \$2}'** Can you explain what does this command do?

**file** shows the file type of a file.

**wc** shows the number of lines, words and characters of a file.

**cat** prints the content of a file in the screen and it also concatenates files.

**head** prints the first 10 lines of a file (or the first **N** lines, with the option **-n <N>**)

**tail** prints the last 10 lines of a file (or the last **N** lines, with the option **-n <N>**)

Can you write a command that prints the **N<sup>th</sup>** line of a text file?

**HINT:** use **head**, **tail** and **|**

## Extra tips

REDACTED

In fact, pipes can be used with almost all commands. They transfer the output of the first command as input for the second one, and so on.

Ex: `ps aux` command  
do?

copy-paste  
copied. To already

`file` shows

`wc` shows

`cat` prints

`head` prints

REDACTED

Check <http://tldp.org/LDP/abs/html/special-chars.html> for a summary of all special characters used in scripts.

Can you write a command that prints the 1st line of a text file?

HINT: use `cat`, `head`, `tail` and `|`

# Outline



- 09:30 – Course starts
- 09:30 – 11:00: Introduction to Linux. Part 1
- 11:00 – 11:15: Coffee break
- 11:15 – 12:30: Introduction to Linux. Part 2
- 12:30 – 14:00: Lunch
- 14:30 – 16:00: Cluster usage. Part 1
- 16:00 – 16:15: Coffee break
- 16:15 – 17:30: Cluster usage. Part 2
- 17:30 – 18:00: Feedback + Q&A



# Clusters

Cluster are HPC infrastructures:

- Parallel computing servers
- Large-memory servers
- Large-storage capacity

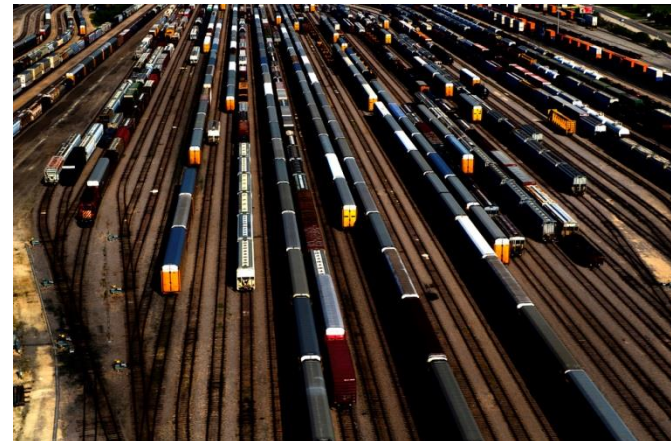


Using HPC facilities does not mean to do the same as in your laptop but in bigger machines!

- Planning
- Knowing your needs
- Being familiar with the code
- Respecting the rules



vs.





# People in sciCORE

Who is who in sciCORE?

[scicore-admin@unibas.ch](mailto:scicore-admin@unibas.ch)

Steering board



Torsten



Michael



Timm

Management & Administration



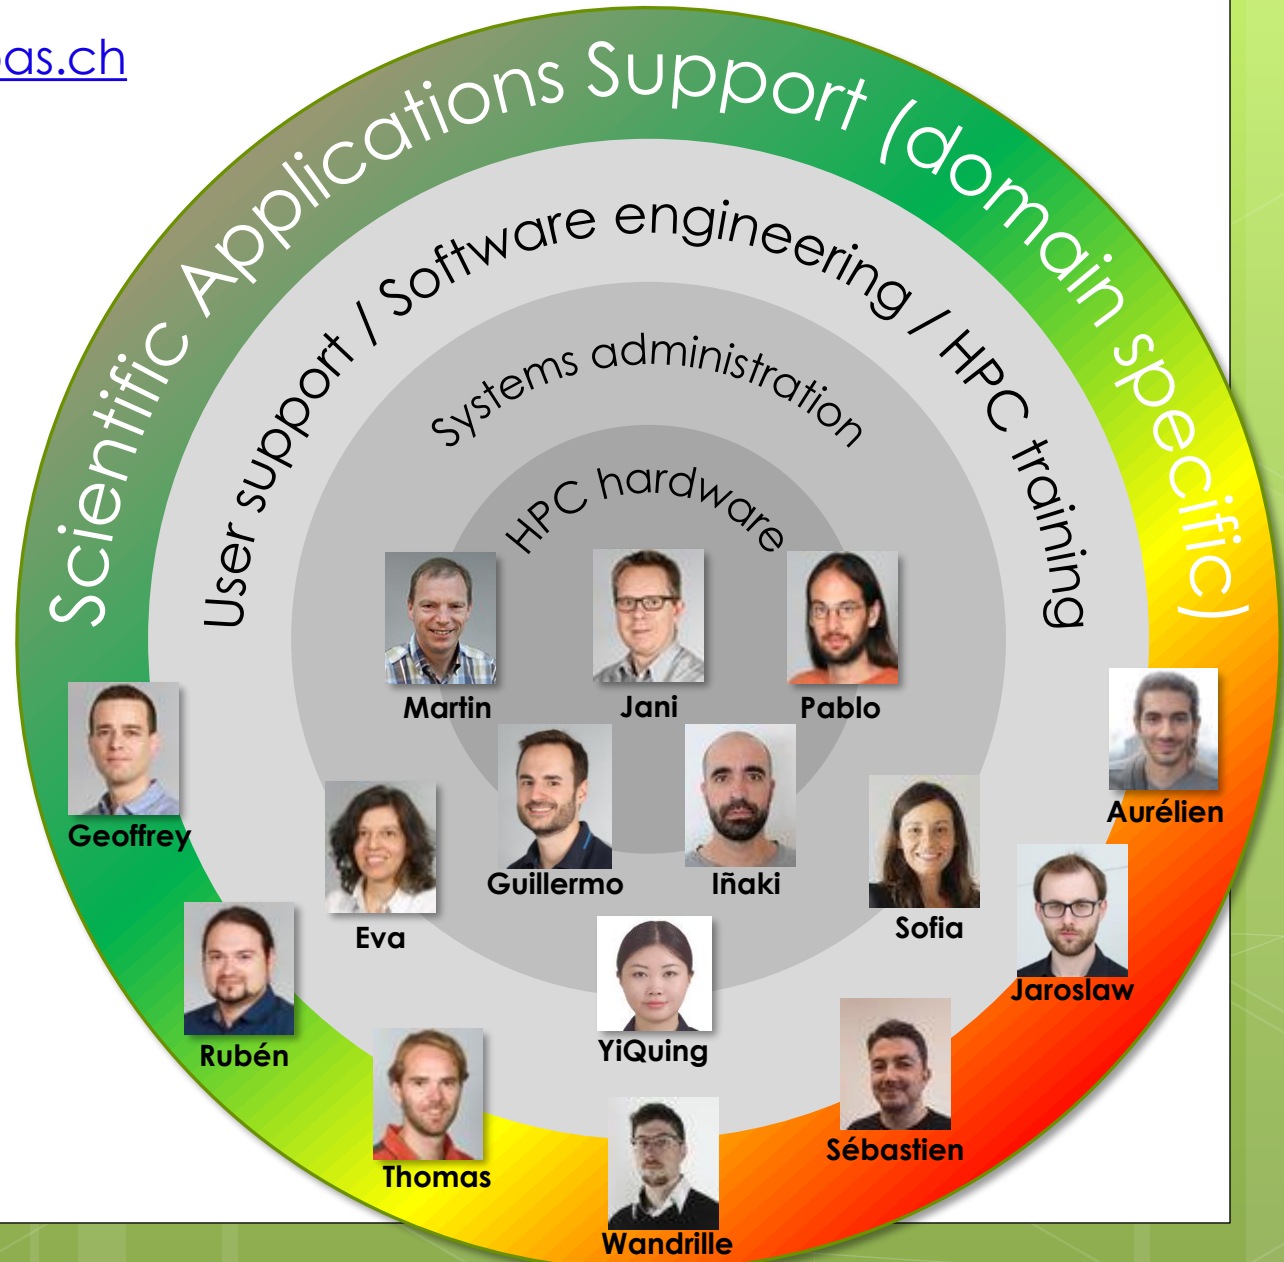
Thierry



Lorenza

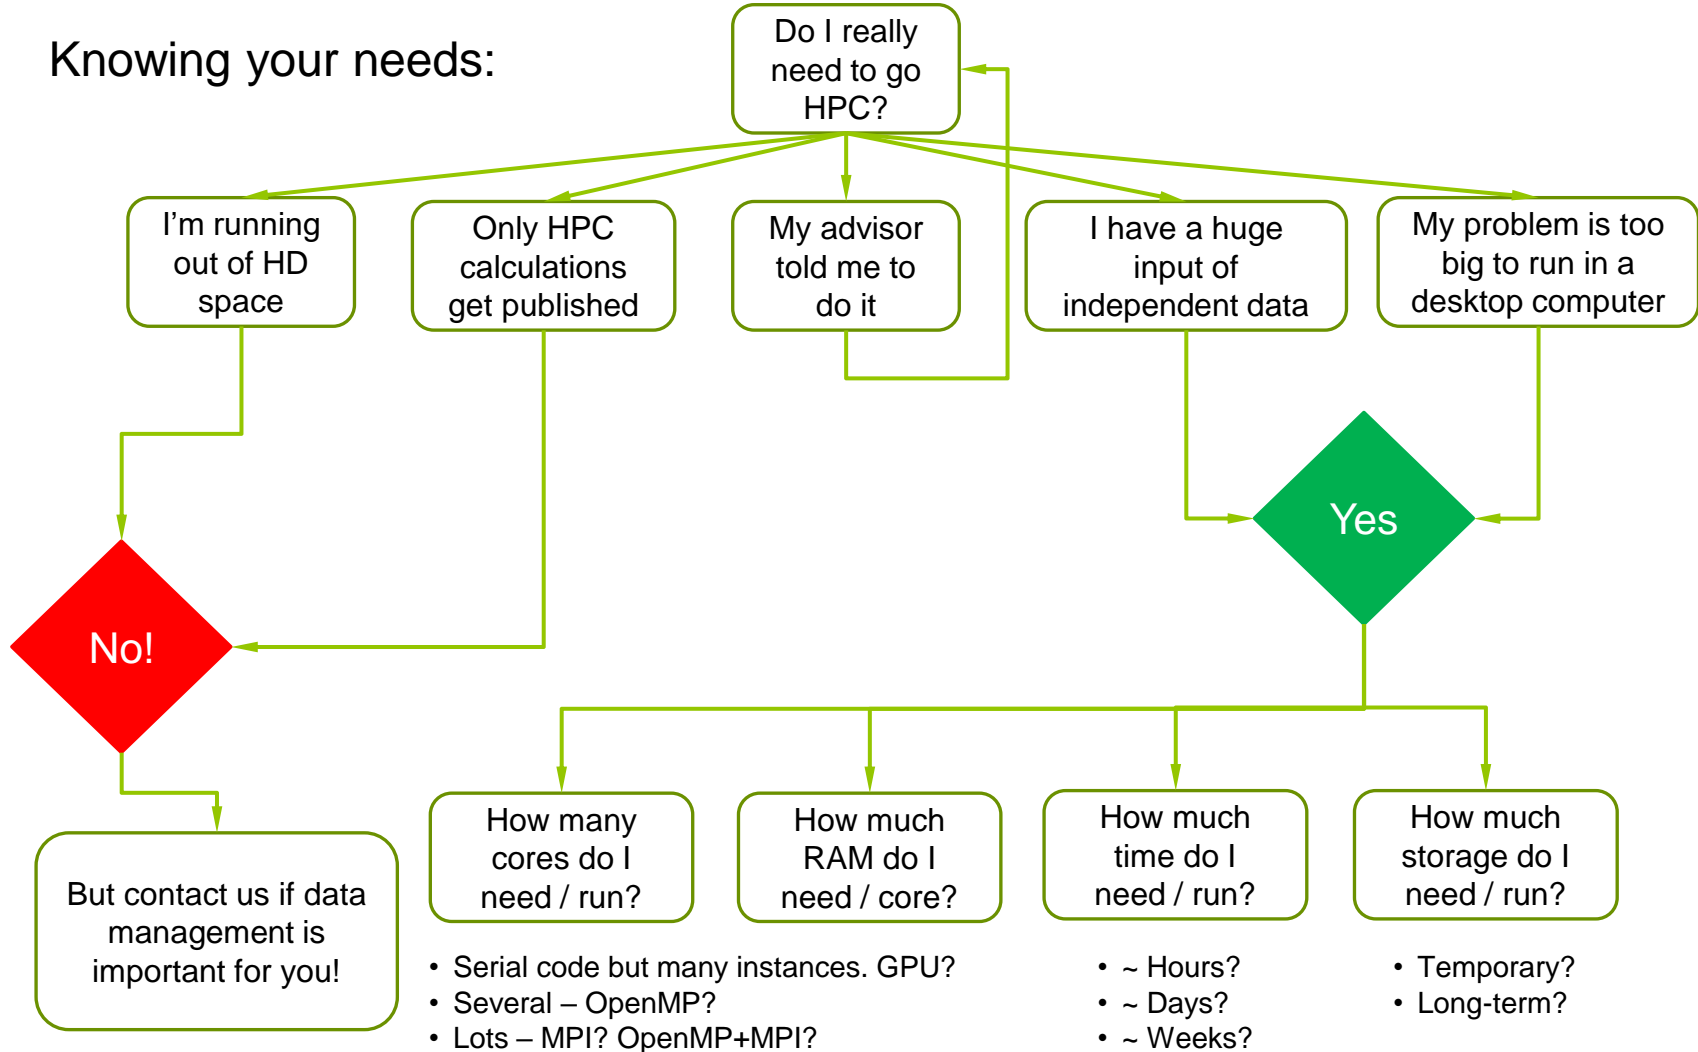


Ariadna



# Using a cluster

## Knowing your needs:



# Using a cluster

Multiple executions of the same code, applied to different sections of the input data



Split the problem in subdomains that are partially solved independently



# Using sciCORE

2 login nodes (front-end to access cluster)

375 computation nodes:

- 8000 cores
- 57 TB RAM (distributed) / 2 node (2 TB) !
- Infiniband interconnected

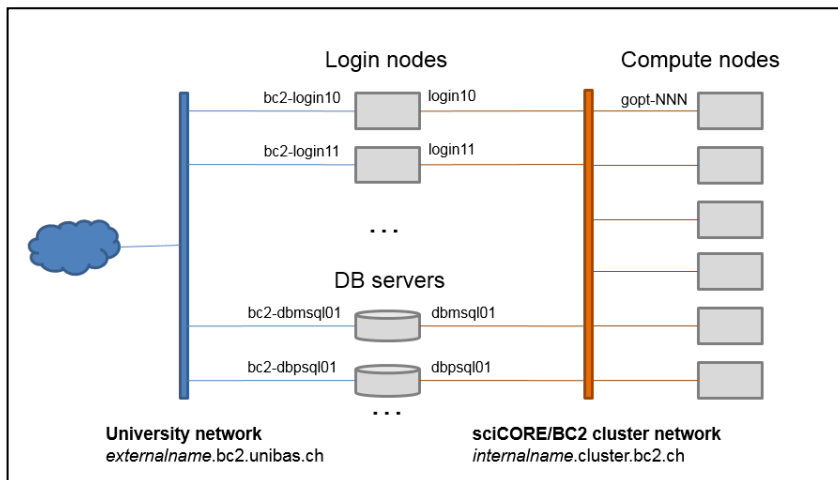
High-performance storage (GPFS):

- 11 PB
- Hierarchical Storage Management (HSM)

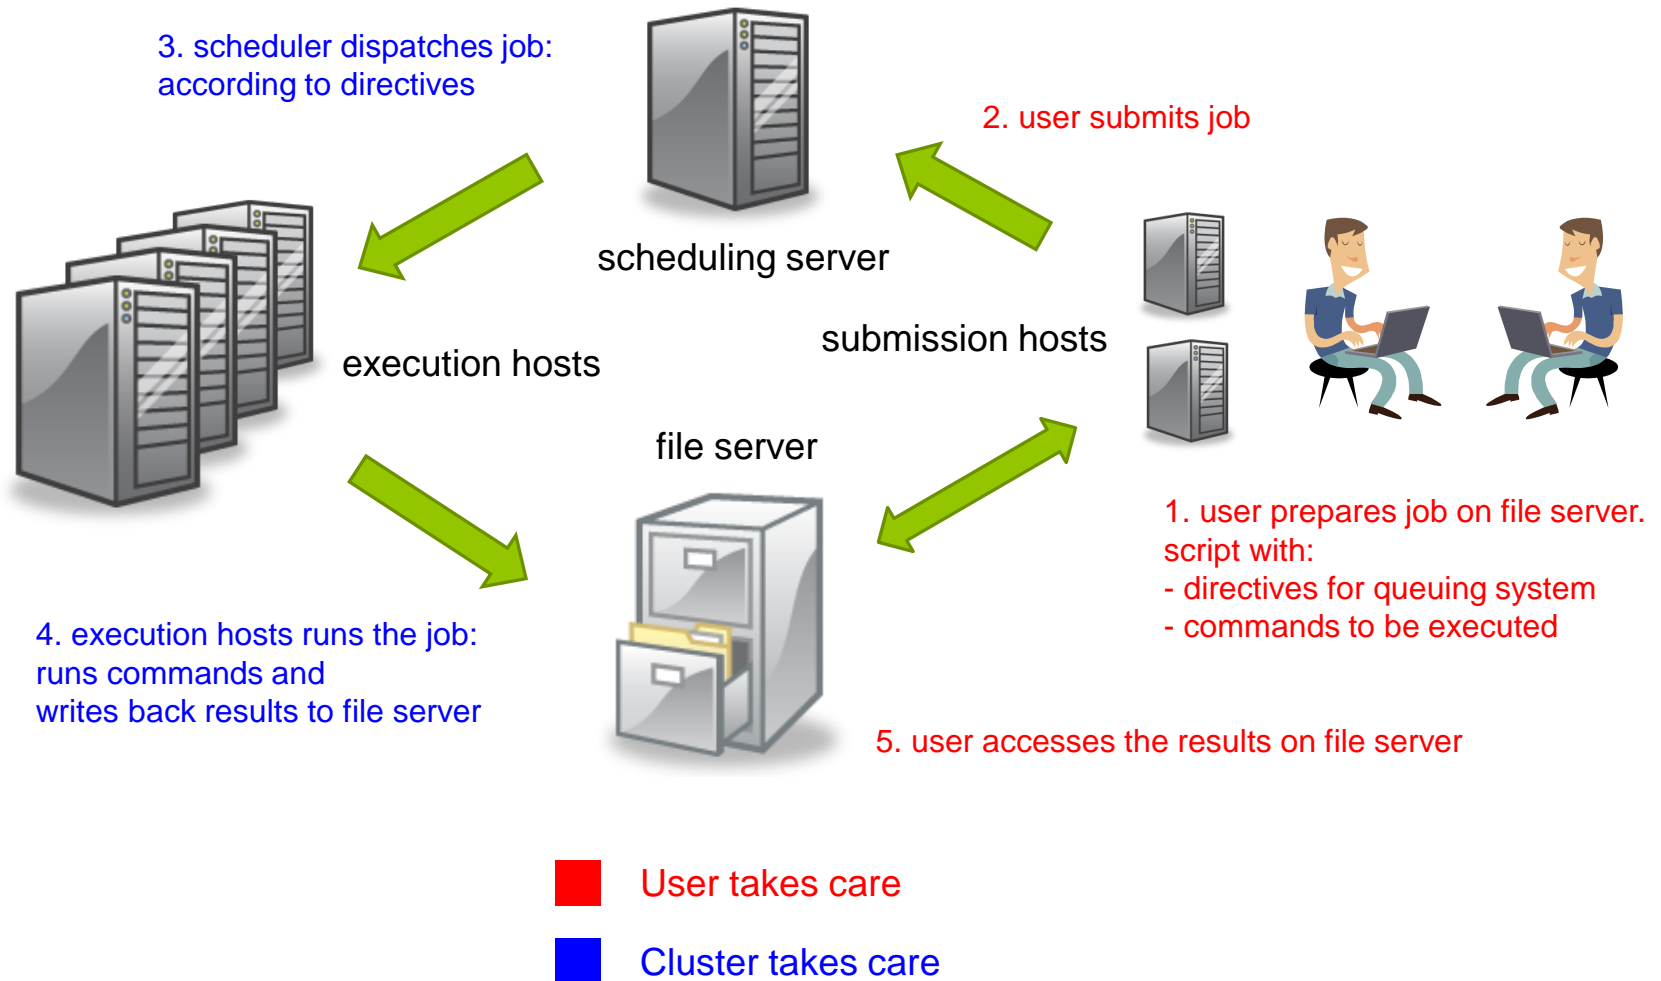
Database and web servers

Private and public services

Purpose	Uni network external name	sciCORE network internalname	Login
Centos 7 – Login node	login	login10 login20	General SLURM
PostgreSQL	bc2-dbpsql01	dbpsql01	No login
MySQL	bc2-dbmsql01	dbmspl01	No login
Web Server	bc2.unibas.ch	www	No login



# Using SLURM



# Using SLURM



## Script for submitting a job

`#!/bin/bash` ← This line is mandatory

```
#SBATCH --job-name=myJob
#SBATCH --cpus-per-task=1
#SBATCH --mem=1G
#SBATCH --time=06:00:00
#SBATCH --qos=6hour
#SBATCH --output=/path/to/stdout/folder
#SBATCH --error=/path/to/stderr/folder
#SBATCH --mail-type=END,FAIL,TIME_LIMIT
#SBATCH --mail-user=<useremail>@unibas.ch
```

All this is optional, but you need to know what is for!

```
# load your required modules
#####
ml Java/1.8.0_92
```

# and here goes your command line

`sleep 30` ← Your command is mandatory

# Using SLURM

## Script for submitting a job

```
#!/bin/bash

#SBATCH --job-name=myJob
#SBATCH --cpus-per-task=1
#SBATCH --mem=1G
#SBATCH --time=06:00:00
#SBATCH --qos=6hour
#SBATCH --output=/path/to/stdout/folder
#SBATCH --error=/path/to/stderr/folder
#SBATCH --mail-type=END,FAIL,TIME_LIMIT
#SBATCH --mail-user=<useremail>@unibas.ch

# load your required modules
#####
ml Java/1.8.0_92

# and here goes your command line
sleep 30
```

Give your job a name that will appear in the queue

Number of processors (default 1 CPU, max: 64)

Total RAM. Can also be requested per CPU with `--mem-per-cpu` (default 1 GB/CPU)

How long will your job run? (default 06:00:00). Scheduler more efficient if this is accurate!

Select queue to run

stdout and stderr (default \$HOME)

Send an email when the specified events occur. Caution!

Define email

Load required software

Execute the program

```
OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK
OPENBLAS_NUM_THREADS=$SLURM_CPUS_PER_TASK
```

Queue name	Max. runtime	Limit TOTAL (cores)	Limit ACCOUNT (cores)	Limit USER (cores)	RAM
30min.q	30 min	1200	1000	400	256 GB/node
6hours.q	6 h	1300	800	300	
1d	1 day	400	60	40	
1week.q	7 days	300	60	40	
2weeks.q	14 days	300	60	40	
infinite.q	∞	100	40	20	

Use command **usage** to know the current status

# Using SLURM

## Script for submitting a job

```
#!/bin/bash

#SBATCH --job-name=myJob
#SBATCH --cpus-per-task=1
#SBATCH --mem=1G
#SBATCH --time=06:00:00
#SBATCH --qos=6hour
#SBATCH --output=/path/to/stdout/folder
#SBATCH --error=/path/to/stderr/folder
#SBATCH --mail-type=END,FAIL,TIME_LIMIT
#SBATCH --mail-user=<useremail>@unibas.ch

# load your required modules
#####
ml Java/1.8.0_92

# and here goes your command line
sleep 30
```

Give your job a name that will appear in the queue

Number of processors (default 1 CPU, max: 64)

Total RAM. Can also be requested per CPU with `--mem-per-cpu` (default 1 GB/CPU)

How long will your job run? (default 06:00:00). Scheduler more efficient if this is accurate!

Select queue to run

stdout and stderr (default \$HOME)

Send an email when the specified events occur. Caution!

Define email

Load required software

Execute the program

```
OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK
OPENBLAS_NUM_THREADS=$SLURM_CPUS_PER_TASK
```

## Working with SLURM

<b>sbatch launch.sh</b>	Sends the script to the queue system, that will schedule the job and assign the required resources.
<b>squeue</b>	Shows the status of the jobs in the queue.
<b>squeue -j &lt;job ID&gt;</b>	Shows detailed data about a specific job.
<b>scancel &lt;job ID&gt;</b>	Cancels a submitted job.
<b>scancel -u &lt;username&gt;</b>	Cancels all submitted jobs owned by <b>&lt;username&gt;</b>
<b>sacct -j &lt;job ID&gt;</b>	Report and account of usage. <b>Ex: sacct -j 25120 -o JobID,AllocCPUS,MaxRSS,State,ExitCode</b> Useful for benchmarking memory needs.
<b>seff &lt;job ID&gt;</b>	Report and account of usage more user friendly. Info given in email by scheduler.



Usually we need additional software to program / compile / debug / launch / process...  
This is handled with a software stack that offers a list of programs/libraries that can be loaded if needed.

[illegible]

## More about modules

Usually we need additional software to program / compile / debug / launch / process...

This is handled with a software stack that offers a list of programs/libraries that can be loaded if needed.

We control what is loaded via the command **ml**

<b>ml</b>	:shows the list of loaded modules.
<b>ml av</b>	:(available) shows the list of all modules.
<b>ml av &lt;keyword&gt;</b>	:shows available modules that contain <keyword> in their name.
<b>ml spider &lt;keyword&gt;</b>	:shows available modules that contain <keyword> in their name with extra info.
<b>ml &lt;module&gt;</b>	:loads <module>
<b>ml -&lt;module&gt;</b>	:unloads <module>
<b>ml swap &lt;module1&gt; &lt;module2&gt;</b>	:unloads <module> and loads <module2> (but better to do purge!)
<b>ml purge</b>	:unloads all modules
<b>ml show &lt;module&gt;</b>	:shows detailed information about <module>

## More about toolchains

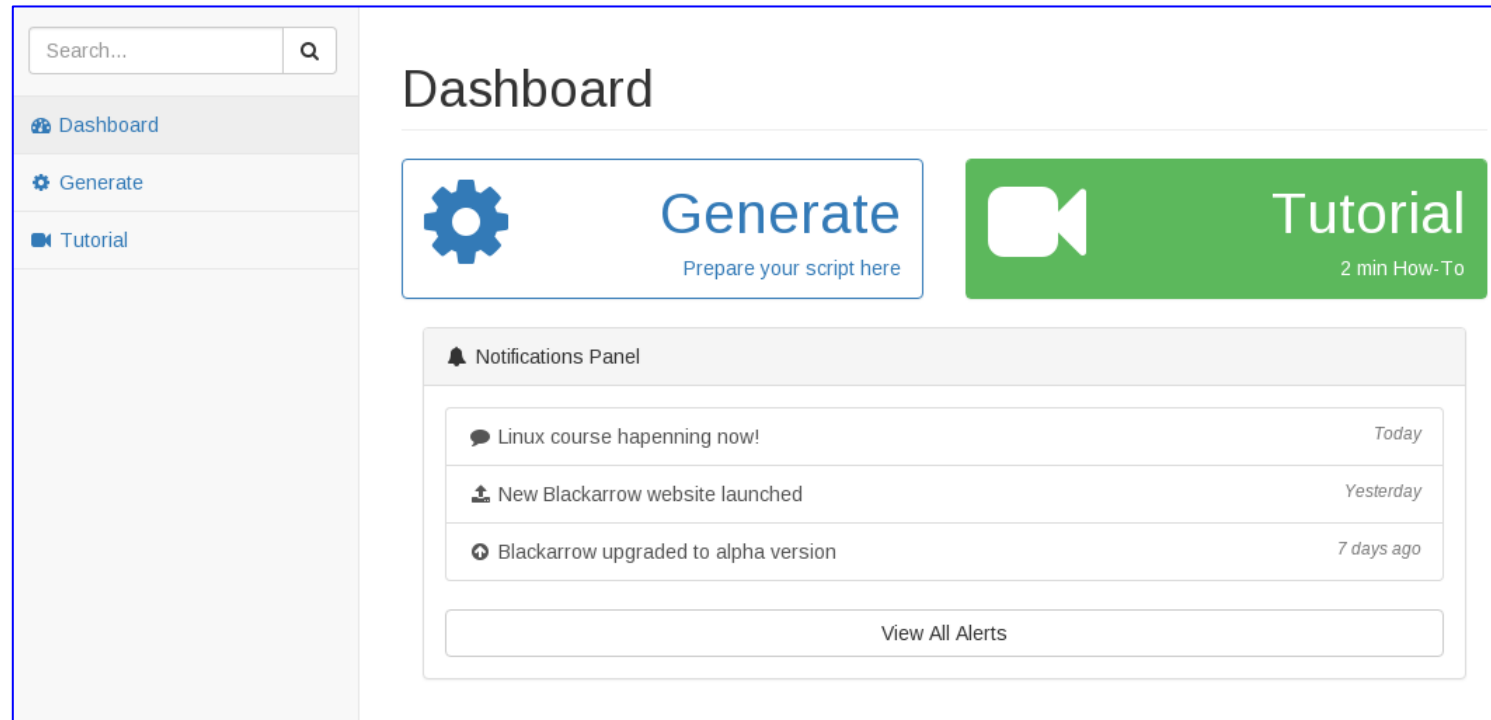
Some software has encoded in its name how it was compiled and with which libraries: toolchains

<b>golf /foss</b>	GNU compiler + OpenMPI + OpenBLAS + FFTW and ScaLAPACK
<b>intel</b>	ifort + icc + MKL + intel MPI
<b>iomkl</b>	ifort + icc + MKL + OpenMPI

You can always do **ml show <keyword>** to know more about it.

# Using SLURM

We designed a user-friendly web-based script generator, that will help you to familiarize with launching scripts.



<https://scriptgen.scicore.unibas.ch>

## 4. Exercise

1. Log in to the cluster
2. Enter in `linux_course/` , create a directory named `Exercise4/` and enter in it.
3. Use scriptgen to write a script named `launch.sh` that contains the following:
4. Save the file and submit it with `sbatch launch.sh`
5. Check the status with `watch squeue -u <username>`
6. Check how long was the code running (`less err.o`)
7. Copy this two files:  
`cp /shared/data/linux/200kheger.relax4 .`  
`cp /shared/data/linux/findneighbors.run .`
9. Additionally, you can copy and read this pdf:  
`cp /shared/data/linux/findneighbors.pdf .`  
to learn a bit more about what does `findneighbors.run` calculate.
10. Create a script to launch `findneighbors.run` for 10 minutes with 300Mb and 1 core.  
And add the commands  
`ulimit -s unlimited`  
`export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK`  
`time ./findneighbors.run` in your script.  
**Tip:** Copy `launch.sh` to `launch2.sh` and modify it correspondingly.
11. Check that the code is running and note the job ID.
12. Get extra info with `sacct -o ALL -j <job ID>`. Try now `seff <job ID>`.
13. Check how long was the code running.
14. Re-launch the program now with 3 cores changing to `--cpus-per-task=3`  
How long did it take now? Check that it indeed used 3 cores.

```
#!/bin/bash
```

```
#SBATCH --job-name=test.<yourname>
#SBATCH --cpus-per-task=1
#SBATCH --mem-per-cpu=1G
#SBATCH --time=00:03:00
#SBATCH --output=out.o
#SBATCH --error=err.o
```

```
time sleep 60
```

## 4. Exercise

1. Log in to the cluster

2. Enter You can check the results of the calculation doing the following:

3. Use s

4. Save

5. Chec

6. Chec

7. Cop

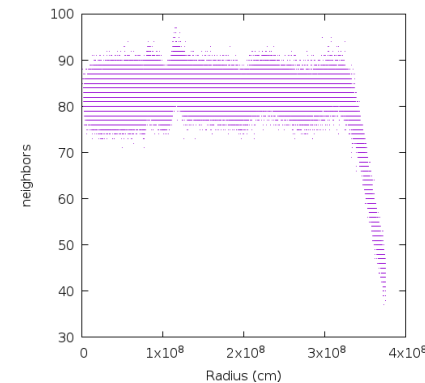
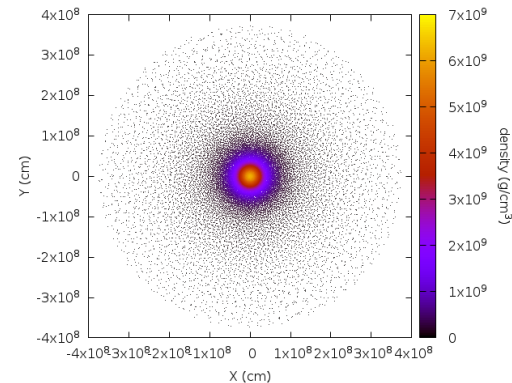
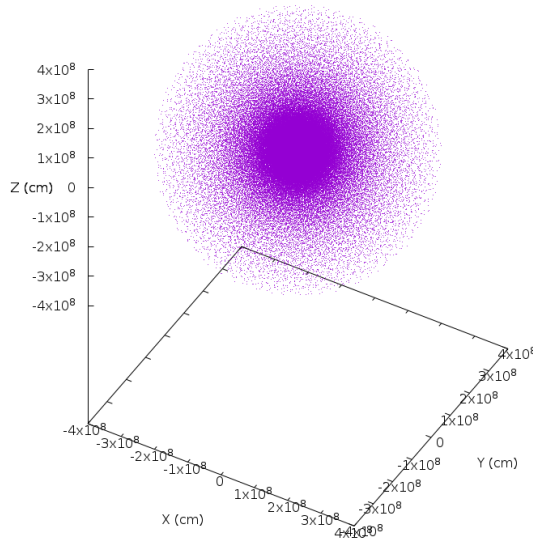
`cp /shared/data/linux/plot.gpl .`

`gnuplot plot.gpl`

(if you get asked to select a module just select #1)

`display results.png`

You should see an image similar to this one:



in it.  
ing:

<yourname>  
1  
1G

1 core.

=3

9. Addi

to le

10. Cre

11. Che

12. Get

13. Che

14. Re-l

How

## That's fine, but...

What happens when I have to submit the same job with different data sets 1,000 times, or 10,000 times?

You should use array jobs:

- You only write one script
- You don't have to worry about deleting thousands of scripts
- If you submit an array job, and realize that you made a mistake, you only have one job id to cancel, instead of 100s.
- You put less burden on the head node.

## Script for submitting an array job

```
#!/bin/bash
```

```
#SBATCH --job-name=myrun
#SBATCH --cpus-per-task=1
#SBATCH --mem-per-cpu=1G
#SBATCH --time=06:00:00
#SBATCH --qos=6hours
#SBATCH --output=myrun.o%A-%3a
#SBATCH --error=myrun.e%A-%3a
#SBATCH --mail-type=END,FAIL,TIME_LIMIT
#SBATCH --mail-user=mailaddress@unibas.ch
```

This is how the output files can be formatted. Here %A stands for the job array's master ID and %a for the job array index. The 3 is for using 3 digits: myrun.o1234-000, myrun.o1234-001, ...

```
# Tell SLURM that this is an array of jobs
#####
#SBATCH --array=1-100%5
```

This will launch 100 tasks to be numbered from 1 to 100.

```
# load your required modules
#####
ml Java
```

Optionally, we can limit the amount of tasks running simultaneously.

```
# and here goes your command line
/bin/sleep 30
```

When a task in the array job is sent to a compute node, its task number is stored in the variable SLURM\_ARRAY\_TASK\_ID, so we can use it to select the input and output commands/data that we want.

## 5. Exercise

```
#!/bin/bash

#SBATCH --job-name=myrun
#SBATCH --cpus-per-task=1
#SBATCH --mem-per-cpu=1G
#SBATCH --time=06:00:00
#SBATCH --output=myrun.o%A-%3a
#SBATCH --error=myrun.e%A-%3a

# Tell SLURM that this is an array of jobs
#####
#SBATCH --array=1-100%10

# and here goes your command line
sleep 30
```



This is just an example to remind you the options of a script.

1. Log in to the cluster
2. Enter in [linux\\_course/](#) and create a directory named [Exercise5/](#) and enter in it.
3. Create a file named [commands.cmd](#) that contains 20 lines with the command [sleep N](#), with [N](#) taking values between 5 and 20.
4. Create a script that launches an array job of 20 jobs in chunks of 5 simultaneous jobs.
5. The script should read the commands from the file [commands.cmd](#) and submit them in the way specified above.
6. When done submit the script and monitorize the submission.

**TIP:** Remember that the variable [\\$SLURM\\_ARRAY\\_TASK\\_ID](#) runs from 1 up the number of jobs in the array.

**TIP:** [\\$\( <command> \)](#) executes [<command>](#)

## Interesting tip:

If you have a pipeline of different processes that depend on the results of the previous step, you can use the option **--dependency**

## Script for submitting an array of dependant jobs

```
#!/bin/bash
```

```
jobID=$(sbatch job1.sh | awk '{print $4}')  
sbatch --dependency=afterok:$jobID job2.sh
```

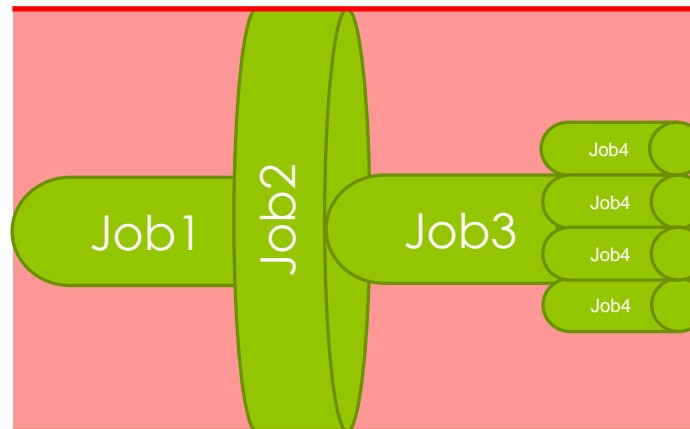
← Get the job ID of the first job of the pipeline.  
← Hold the second job until \$jobID finishes.

```
#!/bin/bash
```

```
sbatch --job-name=myjob1 job1.sh  
sbatch --dependency=singleton --job-name=myjob1 job2.sh
```

← Submit first job.  
← Hold the second job until myjob1 finishes.

Nevertheless, be cautious with very inhomogeneous requirements at different stages of the pipeline.





## Files processing:

**Best practices for Transferring files to/from/on the cluster (and scripts)  
Moving large number of files, and/or large files (and scripts)**

For remote transfer, using **scp** is good enough for most cases:

```
scp <options> <file_origin> <file_destination>
```

You can specify a path in a remote machine as: `<username>@<host>:<path>/<filename>`  
Most typical option is `-rv` (recursive and verbose) to transfer full content of a directory.

### Examples:

```
scp myfile.txt cabazon@login.scicore.unibas.ch:/scicore/home/scicore/cabazon/research/  
scp cabazon@login.scicore.unibas.ch:research/myfile.txt .  
scp -rv cabazon@login.scicore.unibas.ch:research/folder1/ ruben@lapalma1.iac.es:simulation/
```

**NOTE:** scp works over ssh, so you need a working ssh access to the machines involved.

# Using SLURM

## Files processing:

To transfer many and/or large files, using **rsync** is the best option:

```
rsync <options> <file_origin> <file_destination>
```

You can specify a path in a remote machine in the same way as with scp: `<username>@<host>:<path>/<filename>`

Copy/Sync a File on a Local Computer	<code>rsync -zvh backup.tar /tmp/backups/</code>
Copy/Sync a Directory on Local Computer	<code>rsync -avzh /home/rpmpkgs /tmp/backups/</code>
Copy a File from a Local Server to a Remote Server with SSH	<code>rsync -avzhe ssh backup.tar ruben@login.scicore.unibas.ch:/backups/</code>
Show Progress While Transferring Data	<code>rsync -avzhe ssh --progress backup.tar ruben@login.scicore.unibas.ch:/backups/</code>
Automatically Delete source Files after successful Transfer	<code>rsync --remove-source-files -zvh backup.tar /tmp/backups/</code>
Do a Dry Run to test the command before doing any changes	<code>rsync --dry-run --remove-source-files -zvh backup.tar /tmp/backups/</code>

# Using SLURM

## Files processing:

To transfer many and/or large files, using **rsync** is the best option:

```
rsync <options> <file_origin> <file_destination>
```

You can specify a path in a remote machine in the same way as with scp: `<username>@<host>:<path>/<filename>`

### Common options

Copy/Sync a File on a Local C  
Copy/Sync a Directory on Local

**-v**

verbose

Copy a File from a Local Server  
Server with SSH

**-r**

copies data recursively (but don't  
preserve timestamps and permission  
while transferring data

Show Progress While Transfer

**-a**

archive mode allows copying files  
recursively and it also preserves  
symbolic links, file permissions, user &  
group ownerships and timestamps

Automatically Delete source File  
successful Transfer

**-z**

compress file data

Do a Dry Run to test the command  
any changes

**-h**

output numbers in a human-readable  
format


**-e**

Specify the transfer protocol name you  
want to use

### Good praxis!

- Keep your storage low in home/
- Add the string “nobackup” to directories that do not need to be backed up from your home/
- Your programs should not access more than 10s-100s files in a single process.
- When submitting your job, check that it is running as expected with the desired parameters.
- Do not run jobs in the login nodes.
- Read the documentation!

**Thank you  
for your  
attention**



**Resume at:  
14:00**



**Coffee  
break**

